



Small Data GAN

Part 2



友達未來學院

陽明交大 AI 學院 魏澤人
tjw@nctu.edu.tw

About Me

交大 AI 學院/緯創資通 AI 青年講座

Google Developers Expert (Machine Learning)

Organizer, GDG Hualien, Hualien.py, ...

Chief Mathematician, iiNumbers

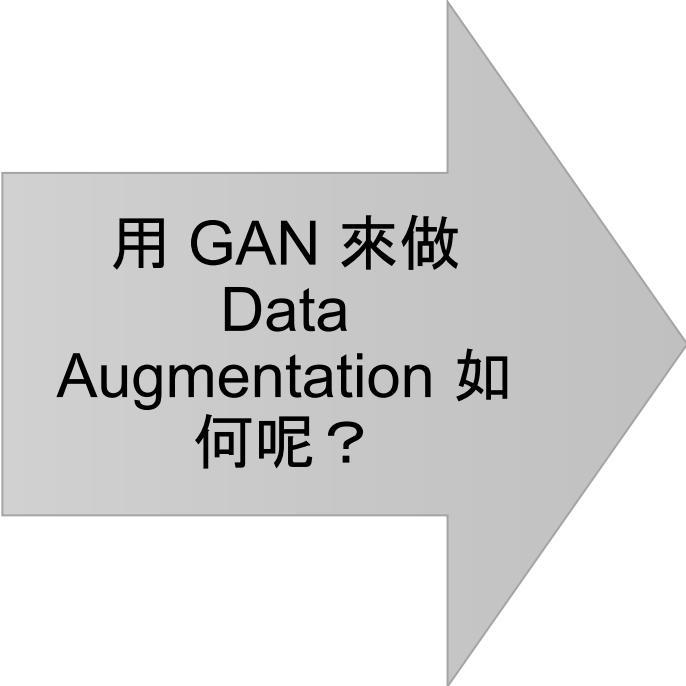
AI and Math Consultant, Ubitus, ITRI, ...

AWS Cloud Ambassador



Review(part 1)

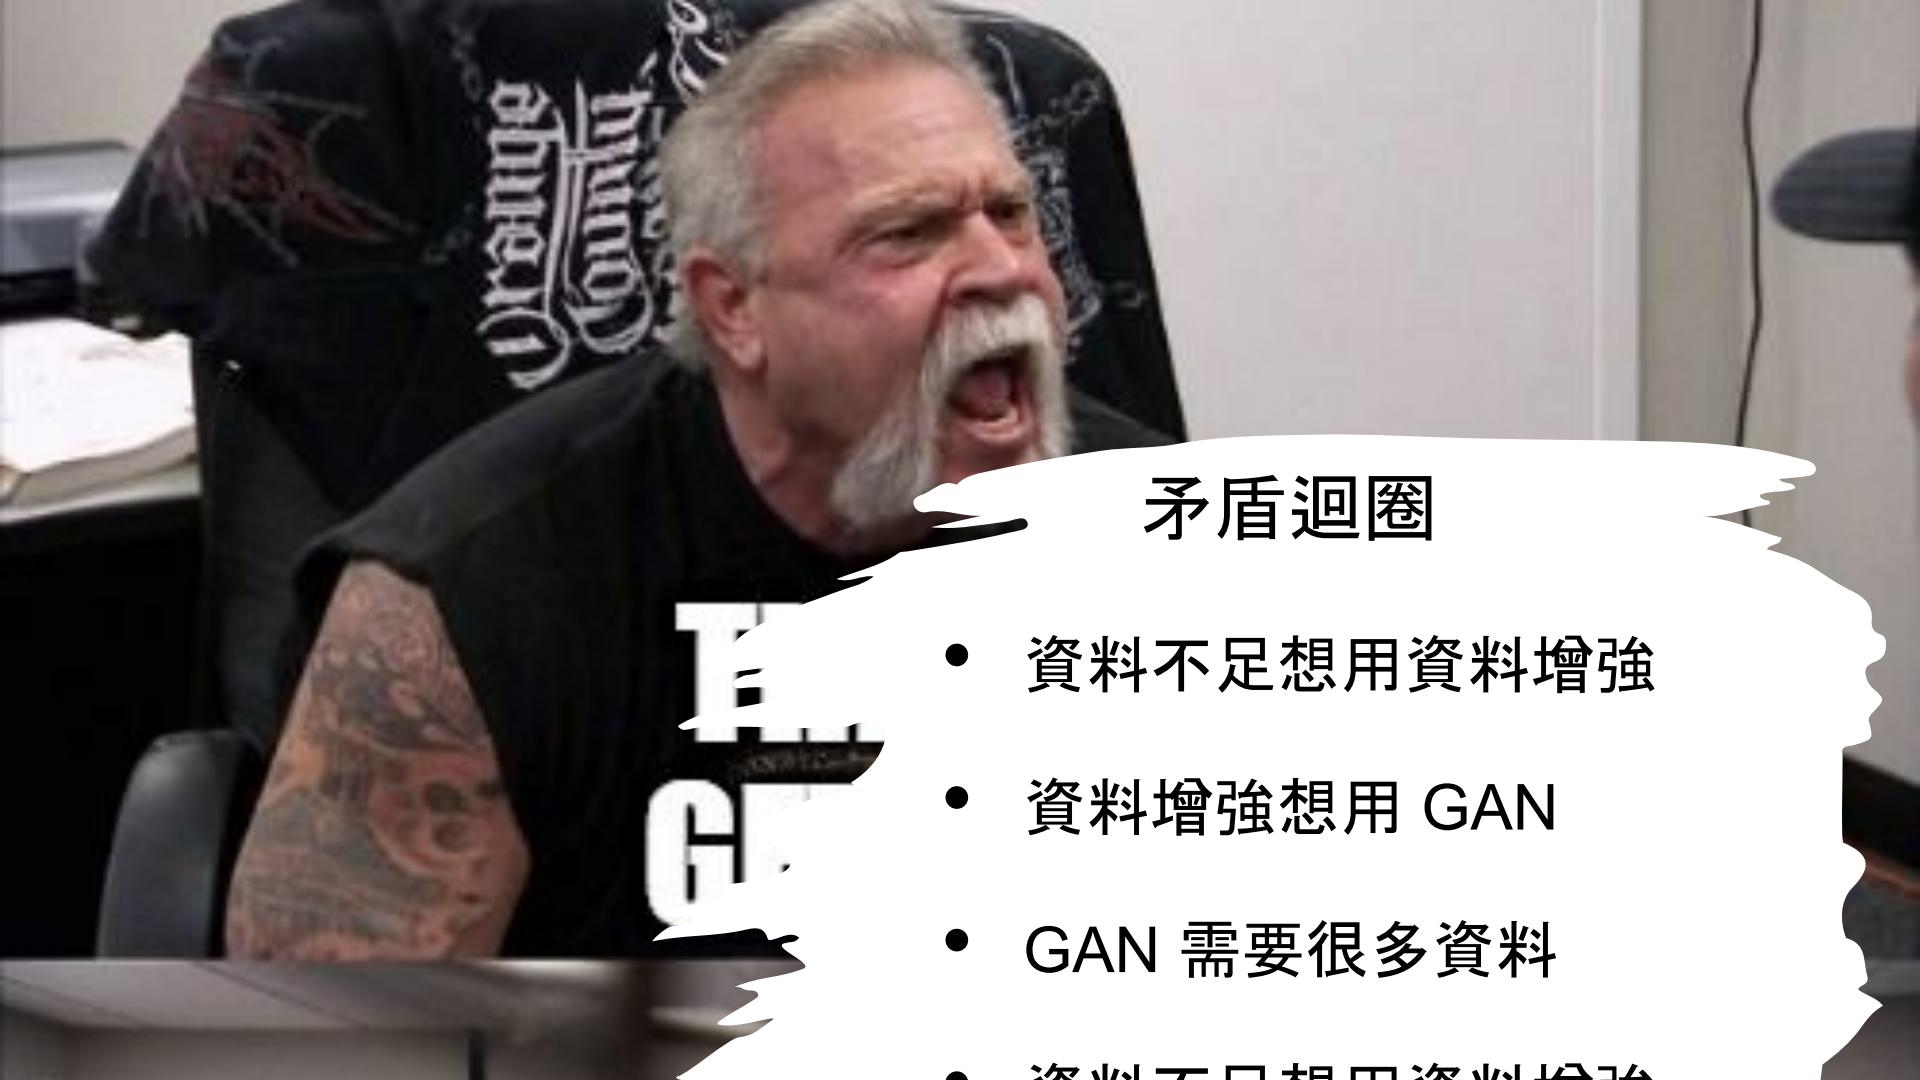
Data Augmentation Using GAN



用 GAN 來做
Data
Augmentation 如
何呢？



情境：資料不足、
資料不平衡



矛盾迴圈

- 資料不足想用資料增強
- 資料增強想用 GAN
- GAN 需要很多資料
- 資料不足想用資料增強

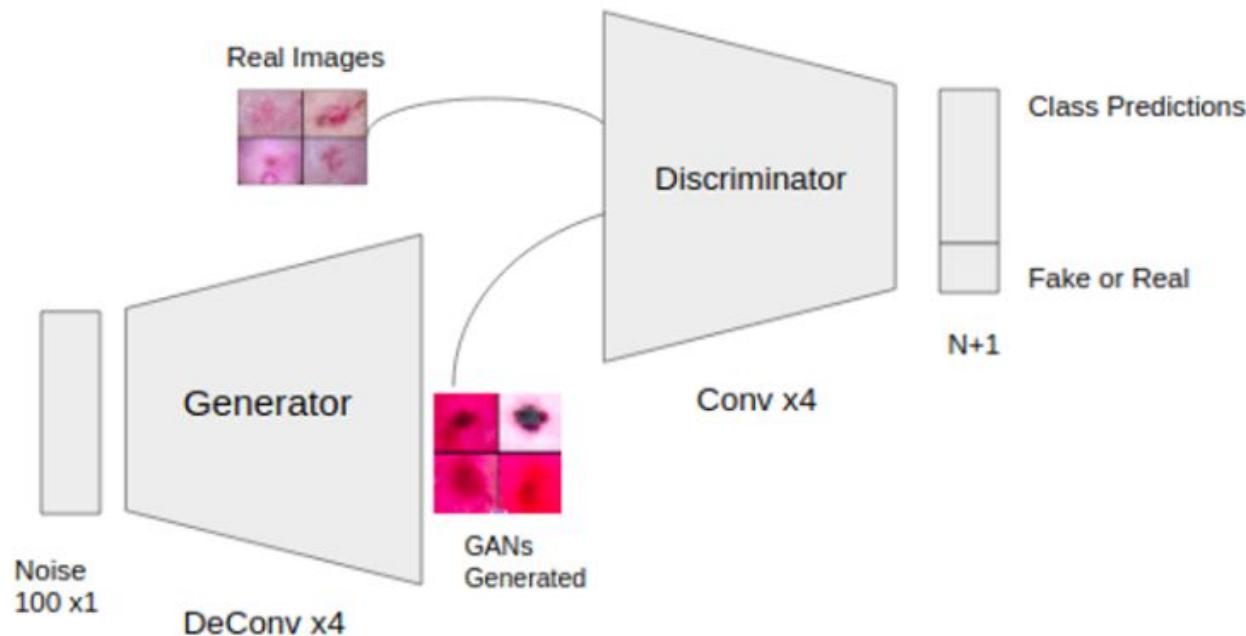
Summary

- Can use classic augmentation to generate training data for GAN.
 - But sometimes this does not work
- Combine classic augmentation with GAN synthesized samples may improve accuracy in some situations.

Summary

- GAN + classic augmentation usually improves the result a little bit
- Some times, this is exactly what you want, the last inch of the last mile.

Architecture: Semi-supervised GAN



Summary

- There are other ways of using GAN to improve the result of other tasks.
- GAN is a way of doing self-supervised learning
- Image to image translation requires fewer training data, compares to unconditional GAN.

AnoGAN

- Schlegl, T., Seeböck, P., Waldstein, S. M., Schmidt-Erfurth, U., Langs, G., 2017a. Unsupervised anomaly detection with generative adversarial networks to guide marker discovery. In: Niethammer, M., Styner, M., Aylward, S., Zhu, H., Oguz, I., Yap, P.-T., Shen, D. (Eds.), Information Processing in Medical Imaging. Springer International Publishing, Cham, pp. 146–157.
- Schlegl, Thomas, et al. "**f-AnoGAN**: Fast unsupervised anomaly detection with generative adversarial networks." *Medical image analysis* 54 (2019): 30-44.
- Shin, Dong-Hoon, Roy C. Park, and Kyungyong Chung. "Decision Boundary-Based Anomaly Detection Model using **Improved AnoGAN** from ECG Data." *IEEE Access* 8 (2020): 108664-108674.
- Zenati, H., Foo, C. S., Lecouat, B., Manek, G., and Chan-drasekhar, V. R. **Efficient GAN-Based Anomaly Detection**. *abs/1802.06222*, 2018. <http://arxiv.org/abs/1802.06222>.
- Di Mattia, Federico, et al. "A survey on gans for anomaly detection." *arXiv preprint arXiv:1906.11632* (2019).
- Akcay, Samet, Amir Atapour-Abarghouei, and Toby P. Breckon. "**Ganomaly**: Semi-supervised anomaly detection via adversarial training." *Asian conference on computer vision*. Springer, Cham, 2018.



Limited Data

How about using
Augmentation to train
GAN?

-
- 當資料量少的時候,
 - 最直接的做法就是 Data augmentation
 - 訓練 GAN 的時候,
 - 如果資料量不夠, 是不是也可以用 Data augmentation?

-
- GAN 是用來學 distribution 的。
 - Augmentation 會改變 data distribution 的特性。
 - 鏡射這些, 可能相對安全,
 - 但旋轉、平移、變色、扭曲, 有可能會改變 distribution.

This is not a No

Improved Consistency Regularization for GANs

Zhengli Zhao^{1,2} Sameer Singh¹ Honglak Lee² Zizhao Zhang² Augustus Odena² Han Zhang²

¹ University of California, Irvine

² Google Research

zhengliz@uci.edu, zhanghan@google.com

Consistency Regularization

- Consistency regularization techniques are popular in the semi-supervised learning literature.
- The basic idea is simple: an input image is **perturbed** in some **semantics-preserving** ways and the **sensitivity** of the **classifier** to that perturbation is **penalized**.

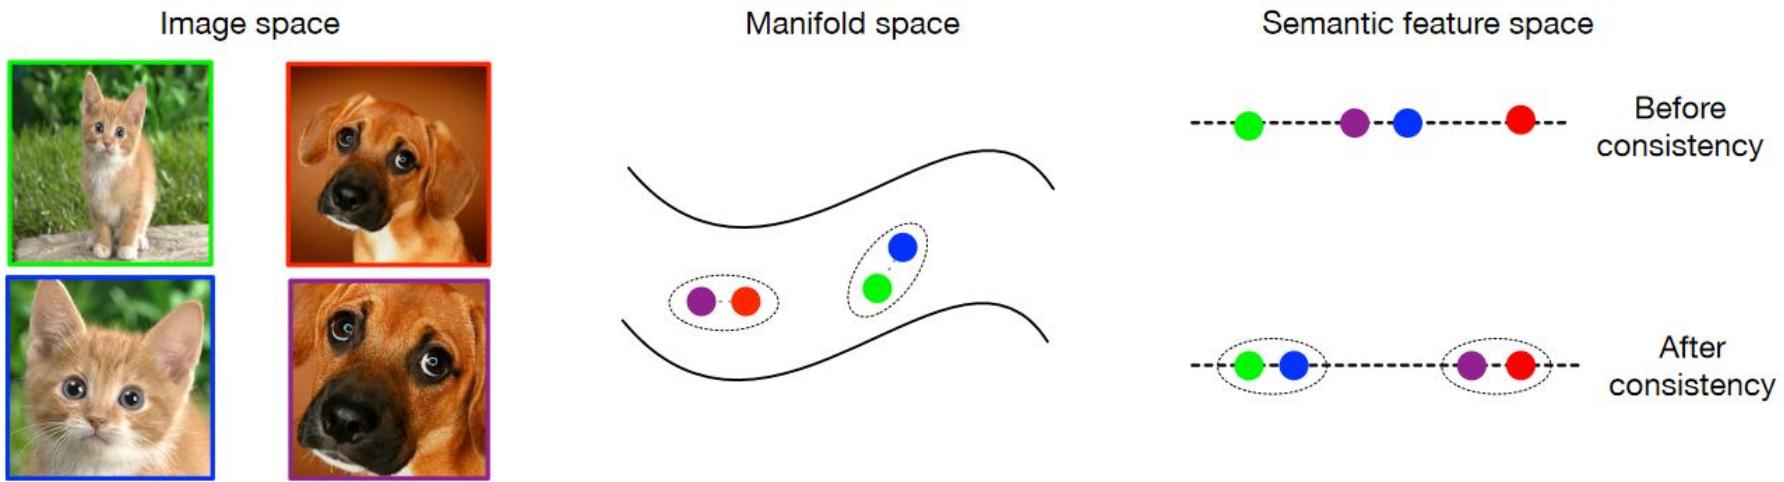
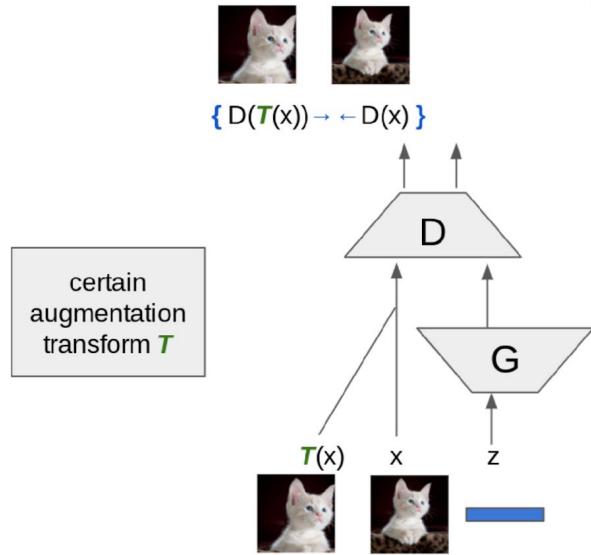


Figure 1: An illustration of consistency regularization for GANs. Before consistency regularization, the zoomed-in dog and the zoomed-in cat (bottom left) can be closer than they are to their original images in feature space induced by the GAN discriminator. This is illustrated in the upper right (the semantic feature space), where the purple dot is closer to the blue dot than to the red dot, and so forth. After we enforce consistency regularization based on the implicit assumption that image augmentation preserves the semantics we care about, the purple dot pulled closer to the red dot.



(1) CR-GAN

$$\min_D L_{cr} = \min_D \sum_{j=m}^n \lambda_j \|D_j(x) - D_j(T(x))\|^2,$$

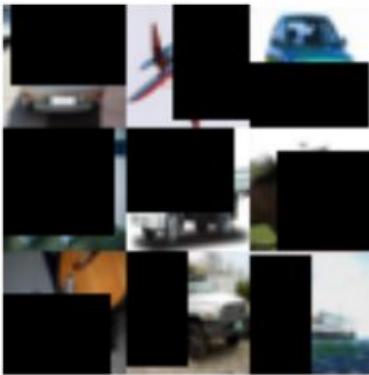
加上這項 regularization 就是 CR-GAN.



(d) 16×16 cutout.



(e) CR samples.



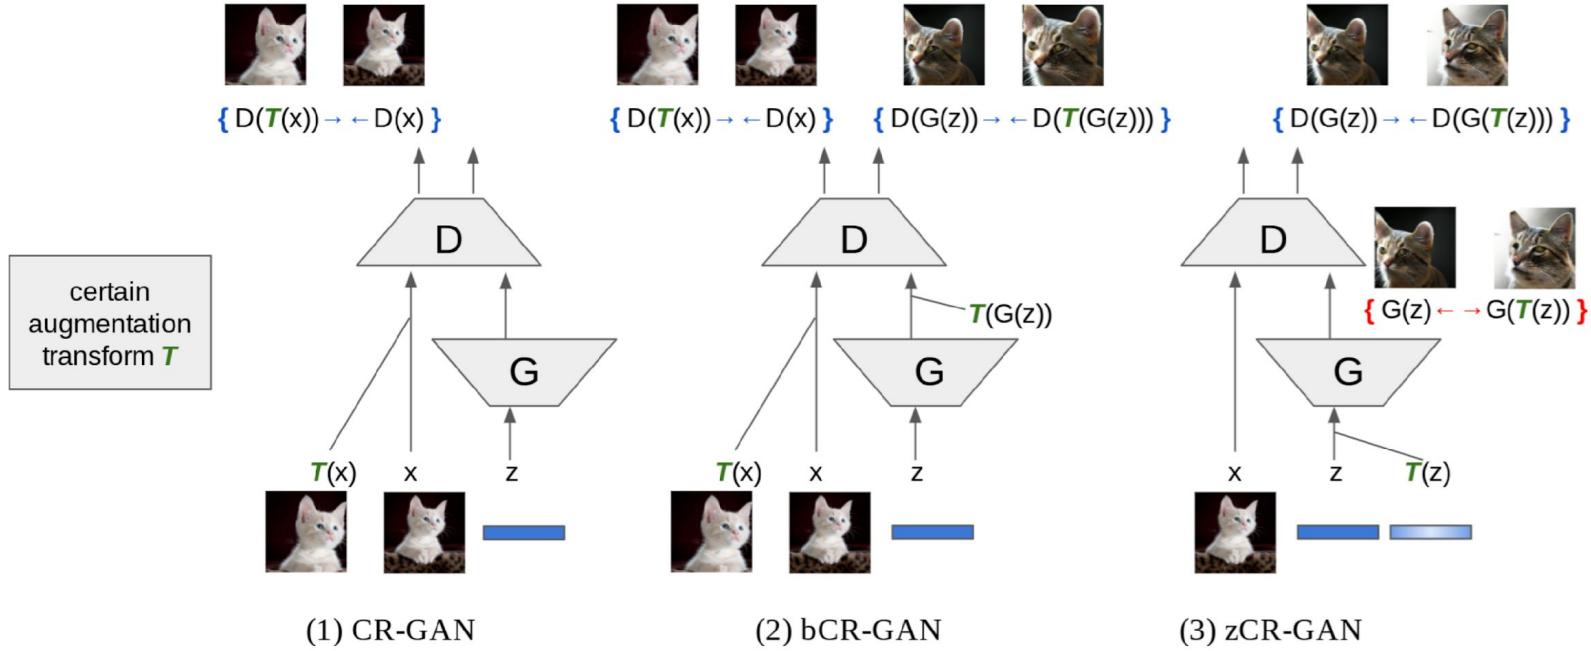
(g) 32×32 cutout.

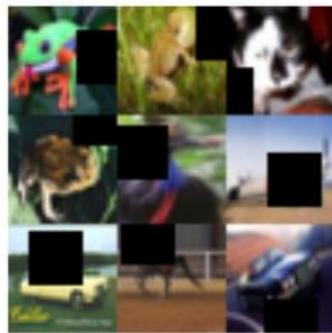


(h) CR samples.

- 但 augmentation 會影響輸出。
- 比方做了 cutout augmentation
- 生出來的圖有時也會產生 cutout。

- CR 是 Consistency Regularization
- bCR 是 balanced 的版本
- zCR 是加上 latent space CR 的版本。

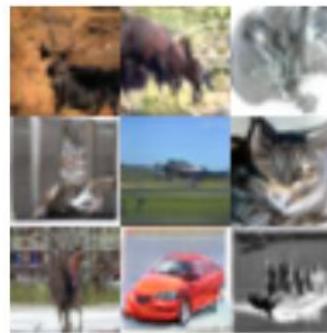




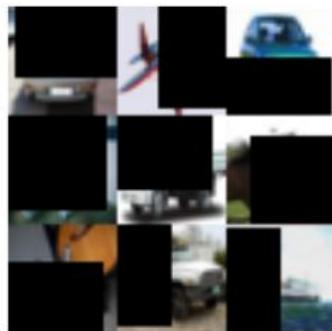
(d) 16×16 cutout.



(e) CR samples.



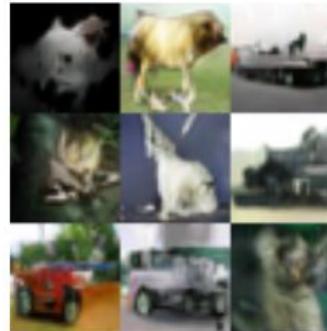
(f) bCR samples.



(g) 32×32 cutout.

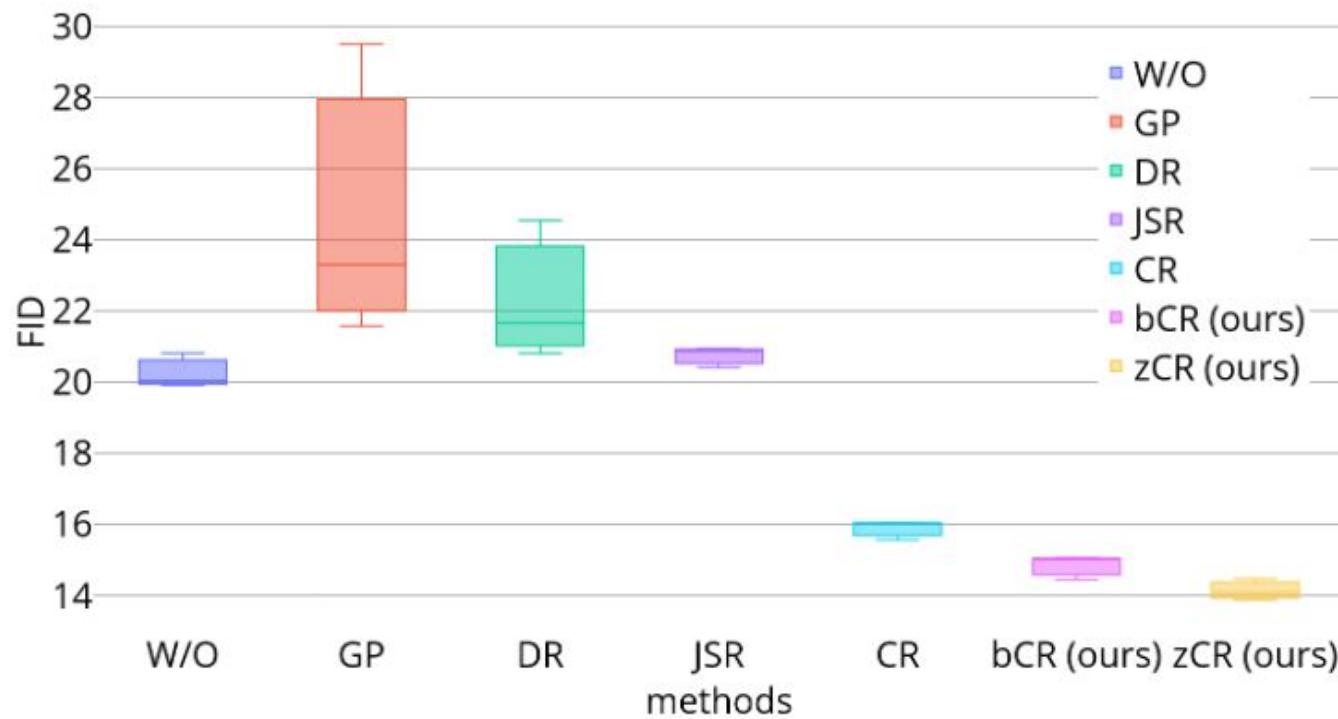


(h) CR samples.



(i) bCR samples.

CIFAR-10, ResNet, non-saturating loss



Models	CIFAR-10	ImageNet
SNGAN	17.50	27.62
BigGAN	14.73	8.73
CR-BigGAN	11.48	6.66
bCR-BigGAN	10.54	6.24
zCR-BigGAN	10.19	5.87
ICR-BigGAN	9.21	5.38

Training Generative Adversarial Networks with Limited Data

Tero Karras
NVIDIA

Miika Aittala
NVIDIA

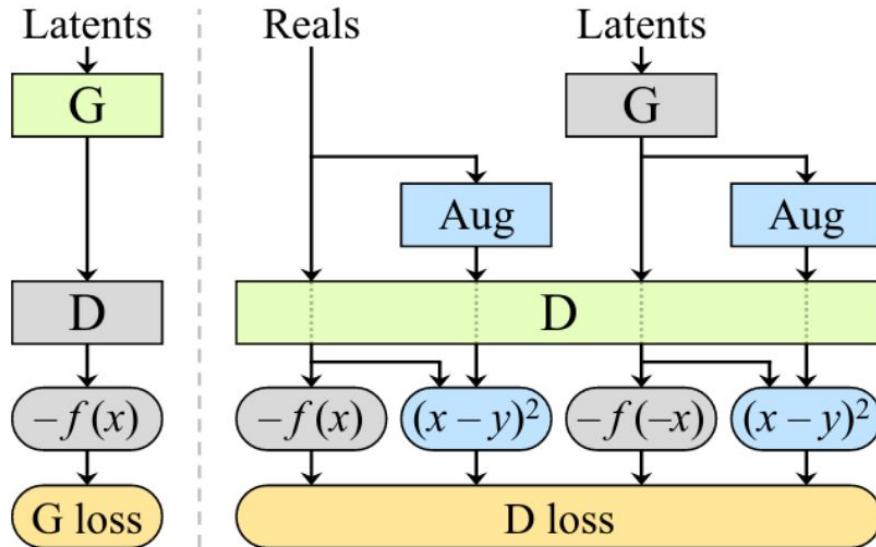
Janne Hellsten
NVIDIA

Samuli Laine
NVIDIA

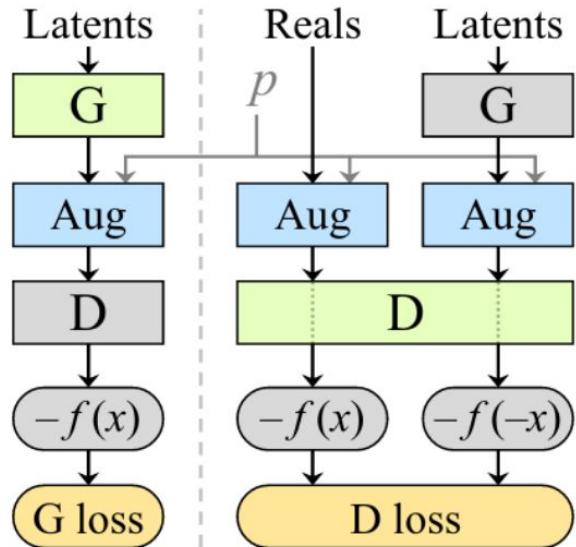
Jaakko Lehtinen
NVIDIA and Aalto University

Timo Aila
NVIDIA

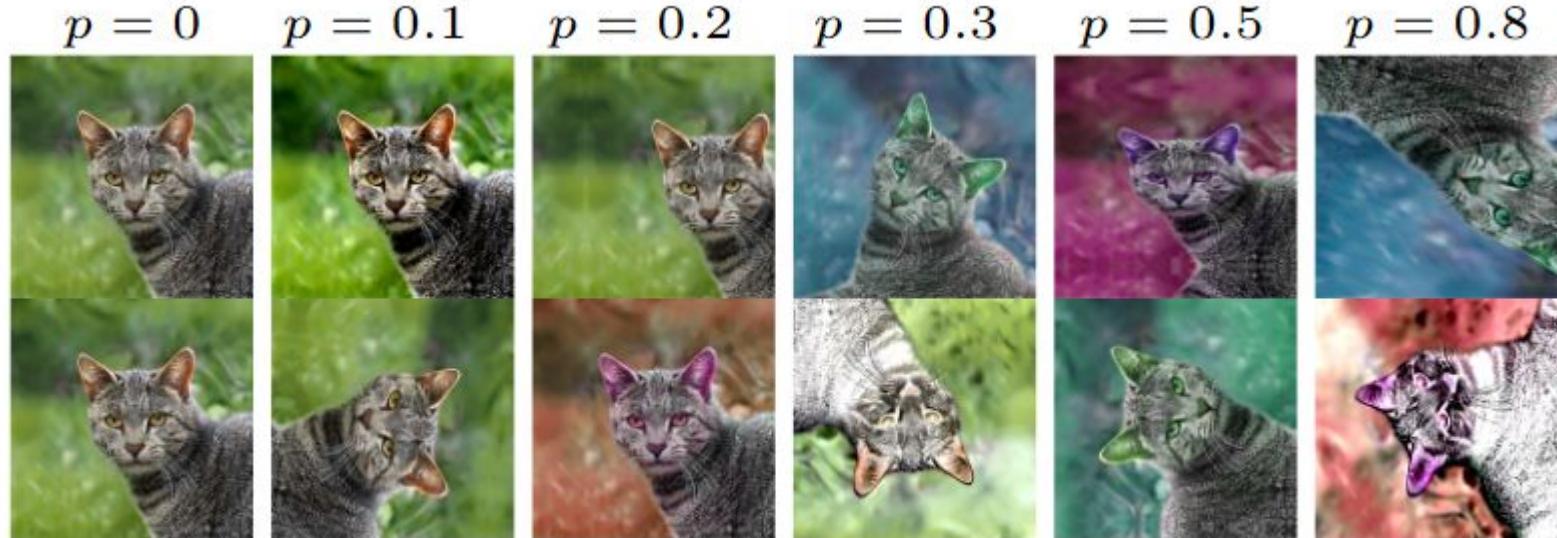
AKA StyleGAN2-ADA



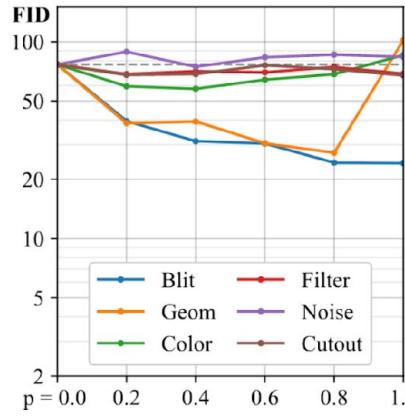
(a) bCR (previous work)



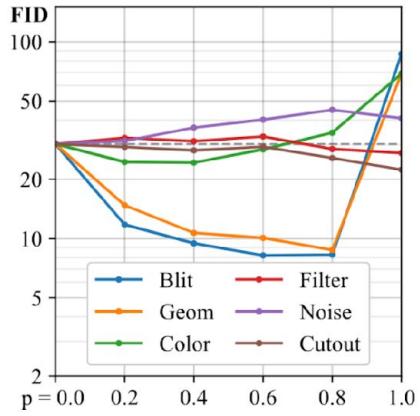
(b) Ours



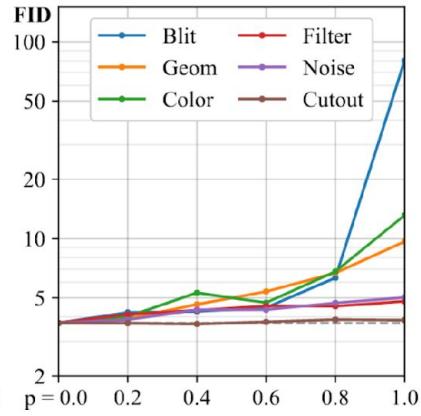
(c) Effect of augmentation probability p



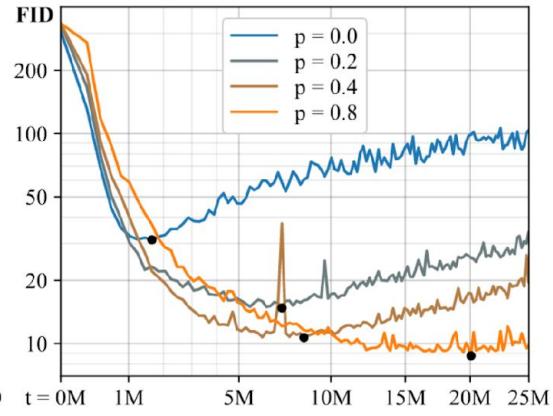
(a) FFHQ-2k



(b) FFHQ-10k



(c) FFHQ-140k



(d) Convergence, 10k, Geom

Figure 4: (a-c) Impact of p for different augmentation categories and dataset sizes. The dashed gray line indicates baseline FID without augmentations. (d) Convergence curves for selected values of p using geometric augmentations with 10k training images.

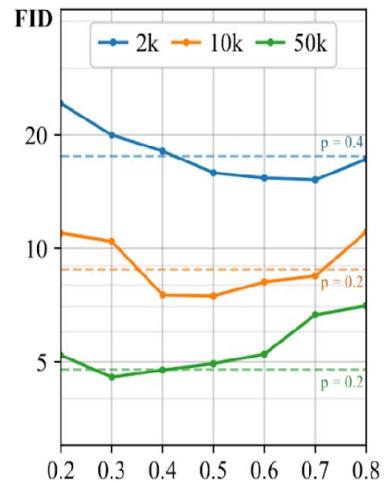
Adaptive discriminator augmentation

Used for preventing overfitting

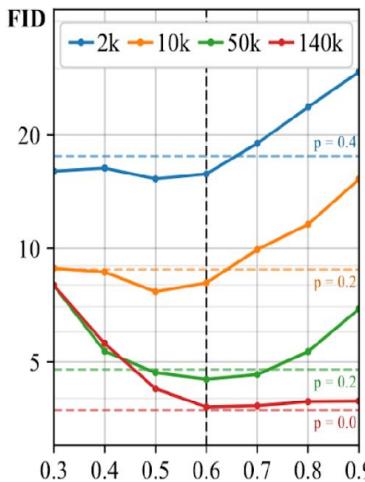
$$r_v = \frac{\mathbb{E}[D_{\text{train}}] - \mathbb{E}[D_{\text{validation}}]}{\mathbb{E}[D_{\text{train}}] - \mathbb{E}[D_{\text{generated}}]}$$

$$r_t = \mathbb{E}[\text{sign}(D_{\text{train}})]$$

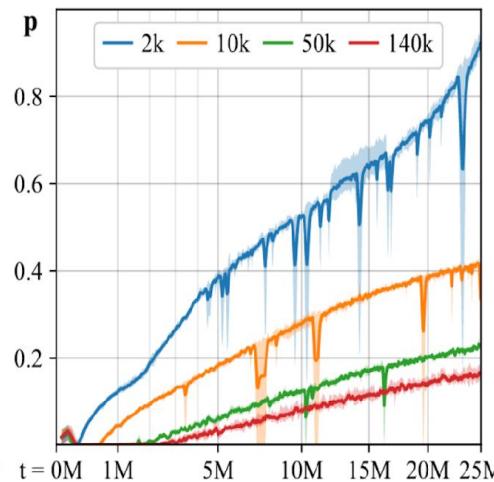
Overfitting Heuristics



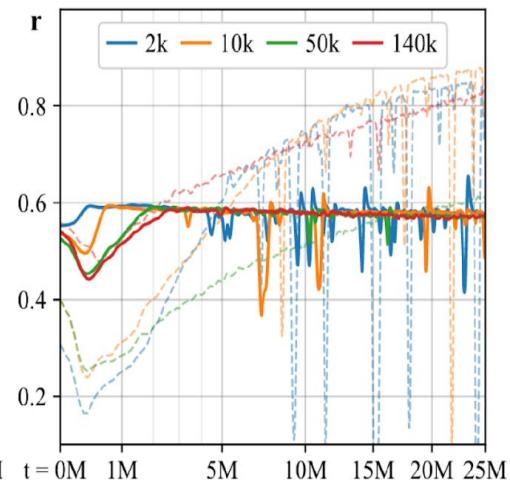
(a) r_v target sweep



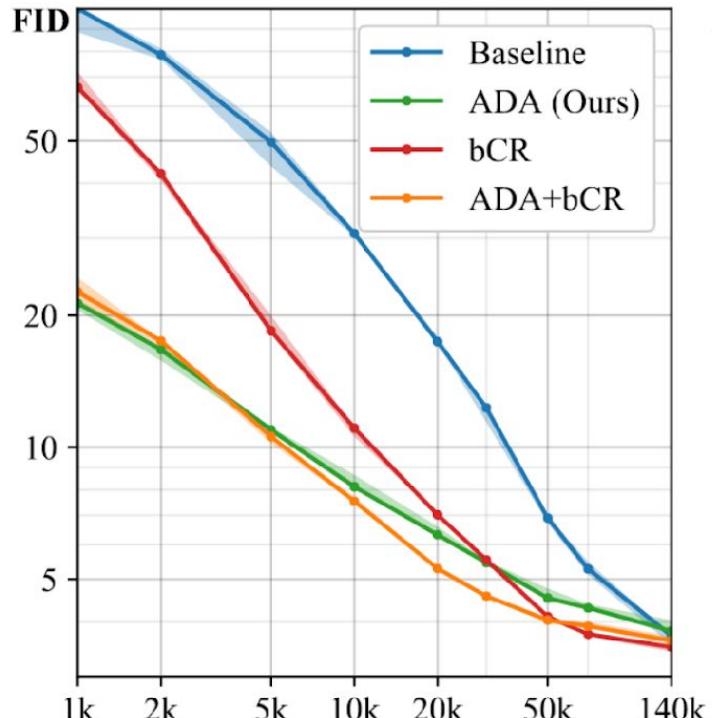
(b) r_t target sweep



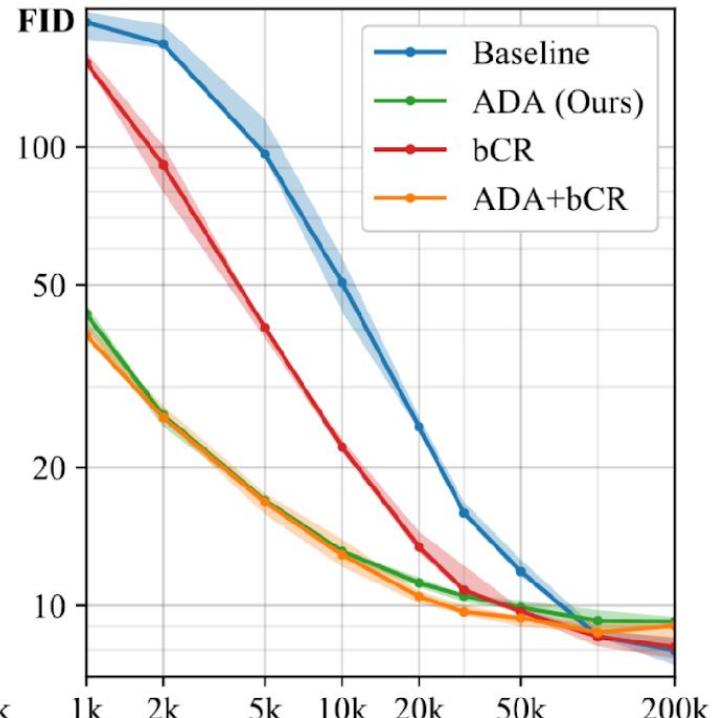
(c) Evolution of p over training



(d) Evolution of r_t

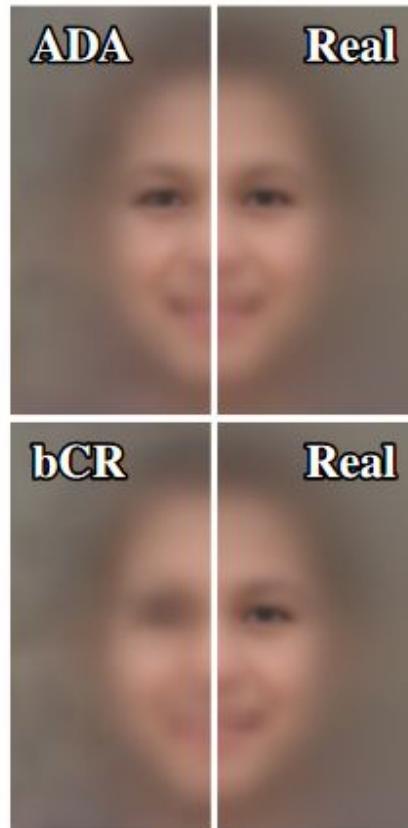


(a) FFHQ (256×256)



(b) LSUN CAT (256×256)

Dataset		Baseline	ADA	+ bCR
FFHQ	1k	100.16	21.29	22.61
	5k	49.68	10.96	10.58
	10k	30.74	8.13	7.53
	30k	12.31	5.46	4.57
	70k	5.28	4.30	3.91
	140k	3.71	3.81	3.62
LSUN CAT	1k	186.91	43.25	38.82
	5k	96.44	16.95	16.80
	10k	50.66	13.13	12.90
	30k	15.90	10.50	9.68
	100k	8.56	9.26	8.73
	200k	7.98	9.22	9.03

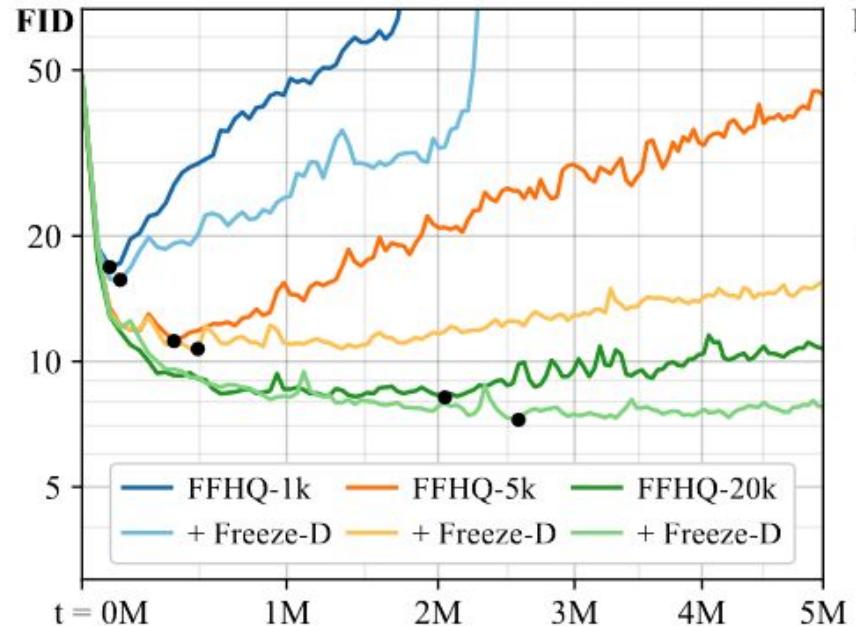


(c) Median FID

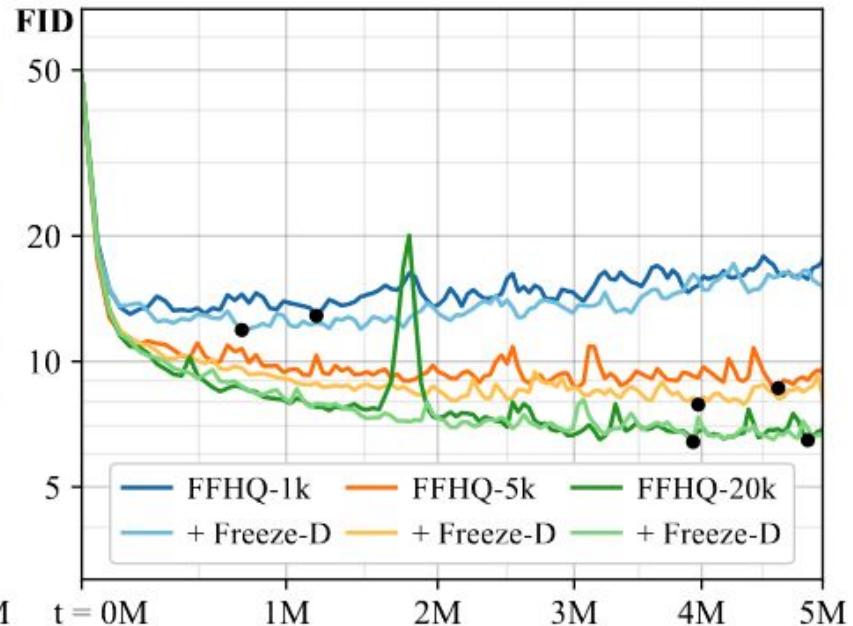
(d) Mean image

FFHQ (256×256)		2k	10k	140k
Baseline		78.80 ± 2.31	30.73 ± 0.48	3.66 ± 0.10
PA-GAN	[48]	56.49 ± 7.28	27.71 ± 2.77	3.78 ± 0.06
WGAN-GP	[15]	79.19 ± 6.30	35.68 ± 1.27	6.54 ± 0.37
zCR	[53]	71.61 ± 9.64	23.02 ± 2.09	3.45 ± 0.19
Auxiliary rotation	[6]	66.64 ± 3.64	25.37 ± 1.45	4.16 ± 0.05
Spectral norm	[31]	88.71 ± 3.18	38.58 ± 3.14	4.60 ± 0.19
Shallow mapping		71.35 ± 7.20	27.71 ± 1.96	3.59 ± 0.22
Adaptive dropout		67.23 ± 4.76	23.33 ± 0.98	4.16 ± 0.05
ADA (Ours)		16.49 ± 0.65	8.29 ± 0.31	3.88 ± 0.13

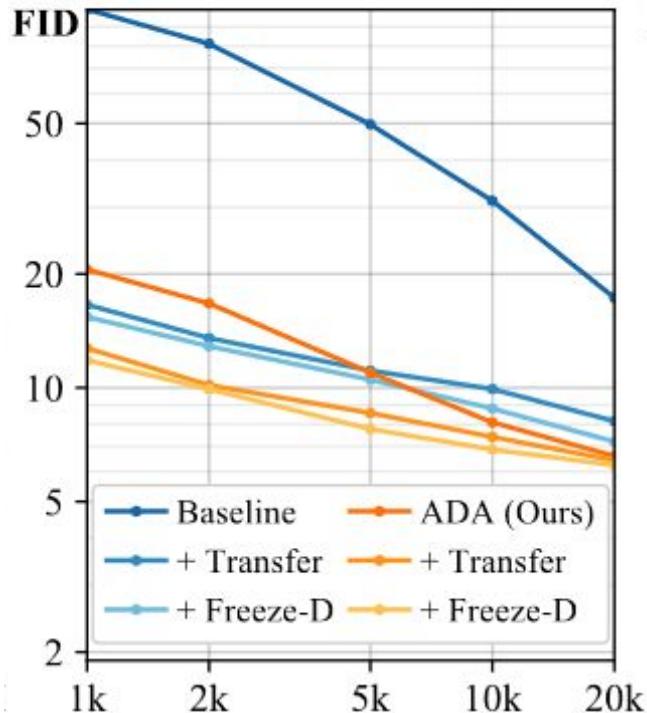
(a) Comparison methods



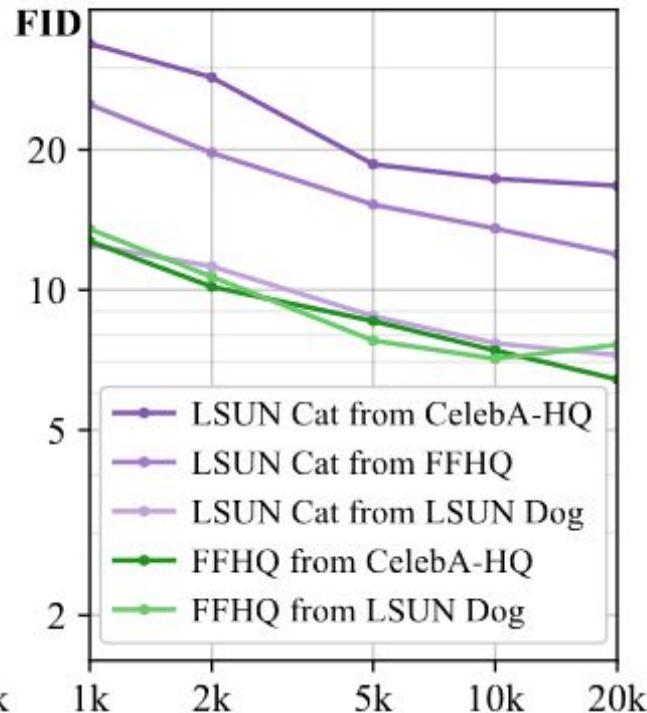
(a) Without ADA



(b) With ADA



(c) Dataset sizes



(d) Datasets



Figure 12: Uncurated 1024×1024 results generated for METFACES (1336 images) with and without ADA, along with real images from the training set. Both generators were trained using transfer learning, starting from the pre-trained StyleGAN2 for FFHQ. We recommend zooming in.

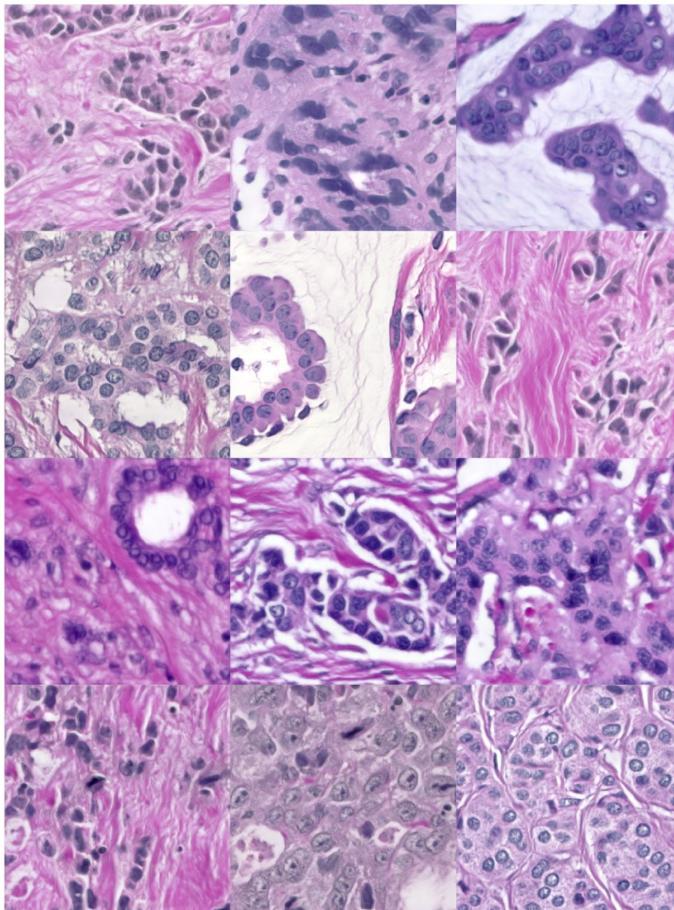
ADA (Ours), untruncated



Original StyleGAN2 config F, untruncated

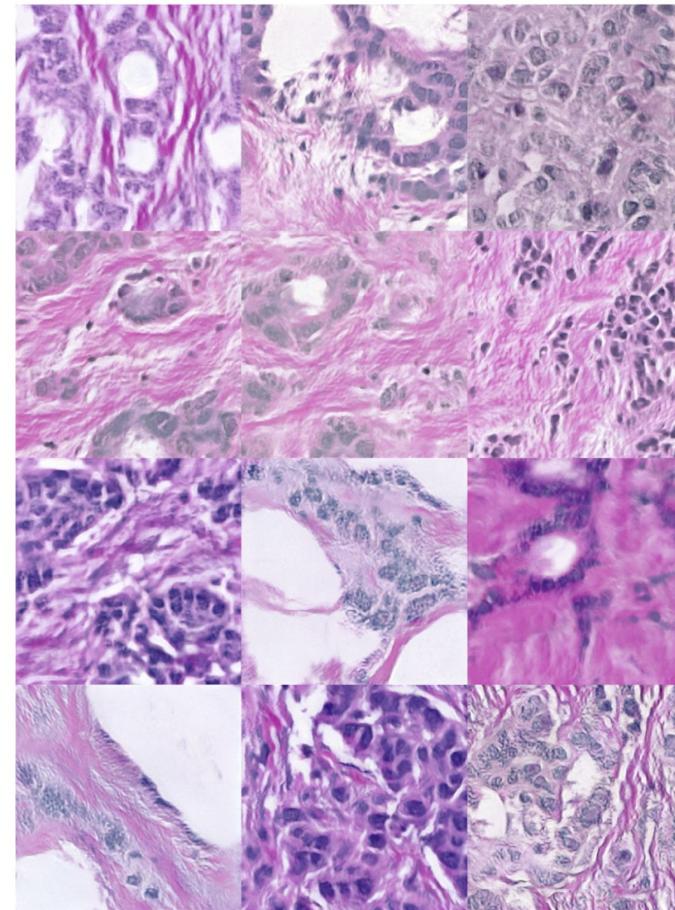
FID 3.55 – KID 0.66×10^3 – Recall 0.430FID 5.13 – KID 1.54×10^3 – Recall 0.215Figure 14: Uncurated 512×512 results generated for AFHQ CAT [7] (5153 images) with and without

ADA (Ours), untruncated



FID 15.71 – KID 2.88×10^3 – Recall 0.340

Original StyleGAN2 config F, untruncated



FID 97.72 – KID 89.76×10^3 – Recall 0.027

Figure 13: Uncurated 512×512 results generated for BRECAHAD [1] (1944 images) with and

Summary

- 用 Augmentation 來訓練 GAN 可能會造成**洩漏**
- 可以控制 Discriminator 看到**平衡的** distribution 來避免洩漏



Few-shot GAN

- <http://nvidia-research-mingyuliu.com/ganimal>
- <https://nvlabs.github.io/FUNIT/>
- Talking head <https://arxiv.org/abs/1905.08233>
- <https://nvlabs.github.io/COCO-FUNIT/paper.pdf>



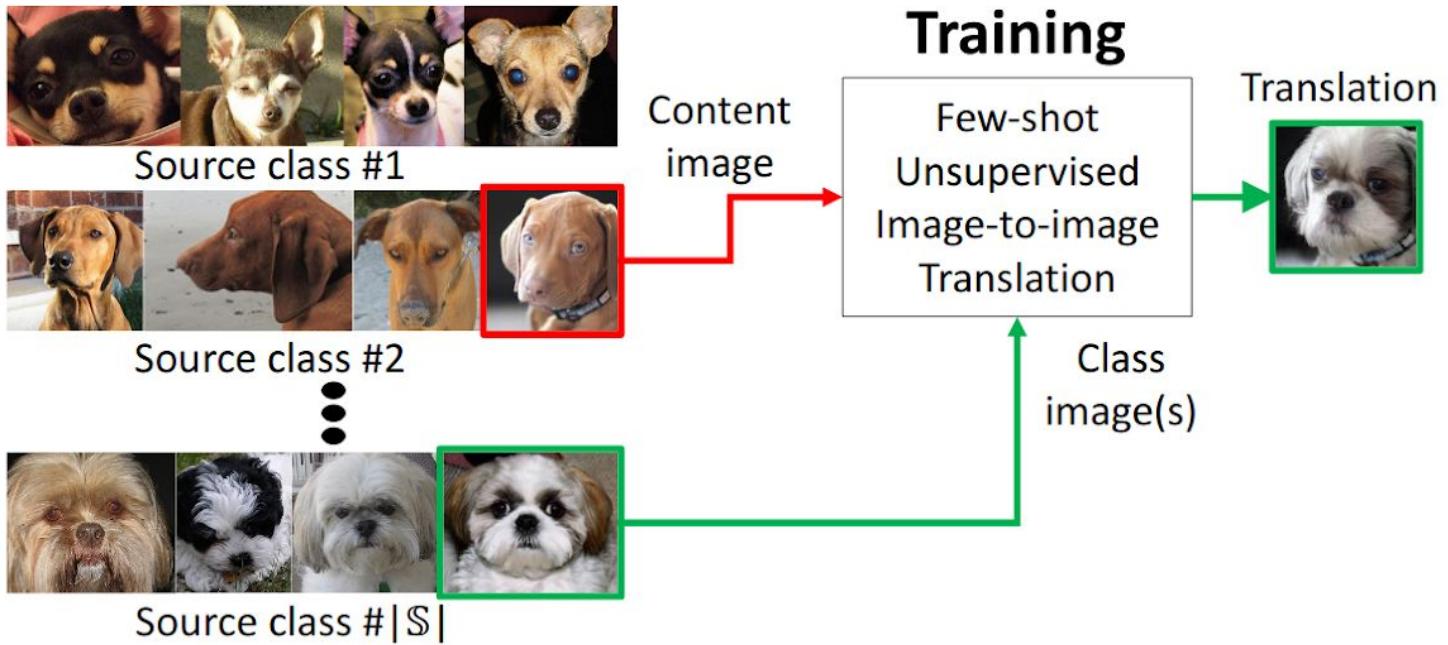
- 1. Input**
- 2. Cocker Spaniel**
- 3. Appenzeller**
- 4. Rottweiler**
- 5. Komondor**
- 6. Kerry Blue Terrier**
- 7. Lakeland Terrier**
- 8. Greater Swiss Mountain Dog**
- 9. Lynx**
- 10. Pug**
- 11. Grey fox**
- 12. Affenpinscher**
- 13. Egyptian Cat**
- 14. Cougar**
- 15. Toy Poodle**
- 16. Red Fox**

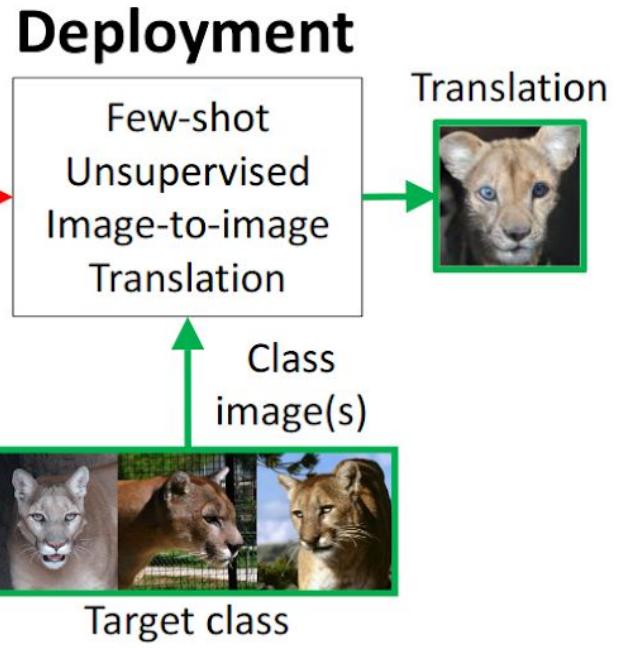
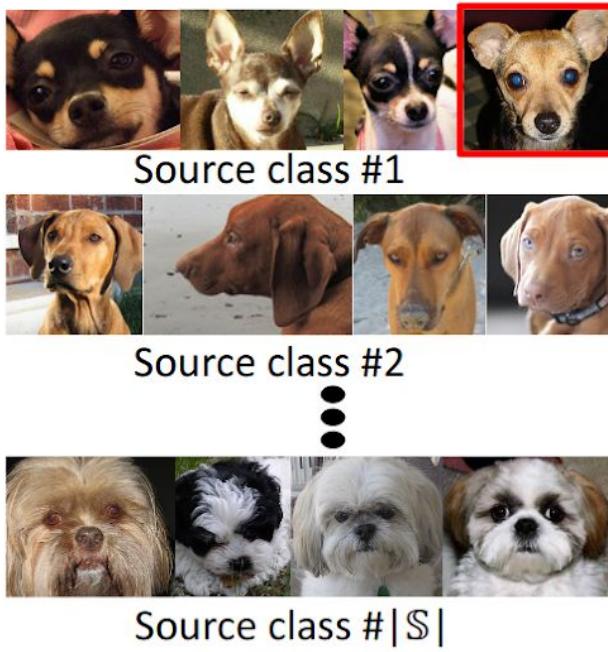
Few-Shot Unsupervised Image-to-Image Translation

Ming-Yu Liu¹, Xun Huang^{1,2}, Arun Mallya¹, Tero Karras¹, Timo Aila¹, Jaakko Lehtinen^{1,3}, Jan Kautz¹

¹NVIDIA, ²Cornell University, ³Aalto University

{mingyul, xunh, amallya, tkarras, taila, jlehtinen, jkautz}@nvidia.com

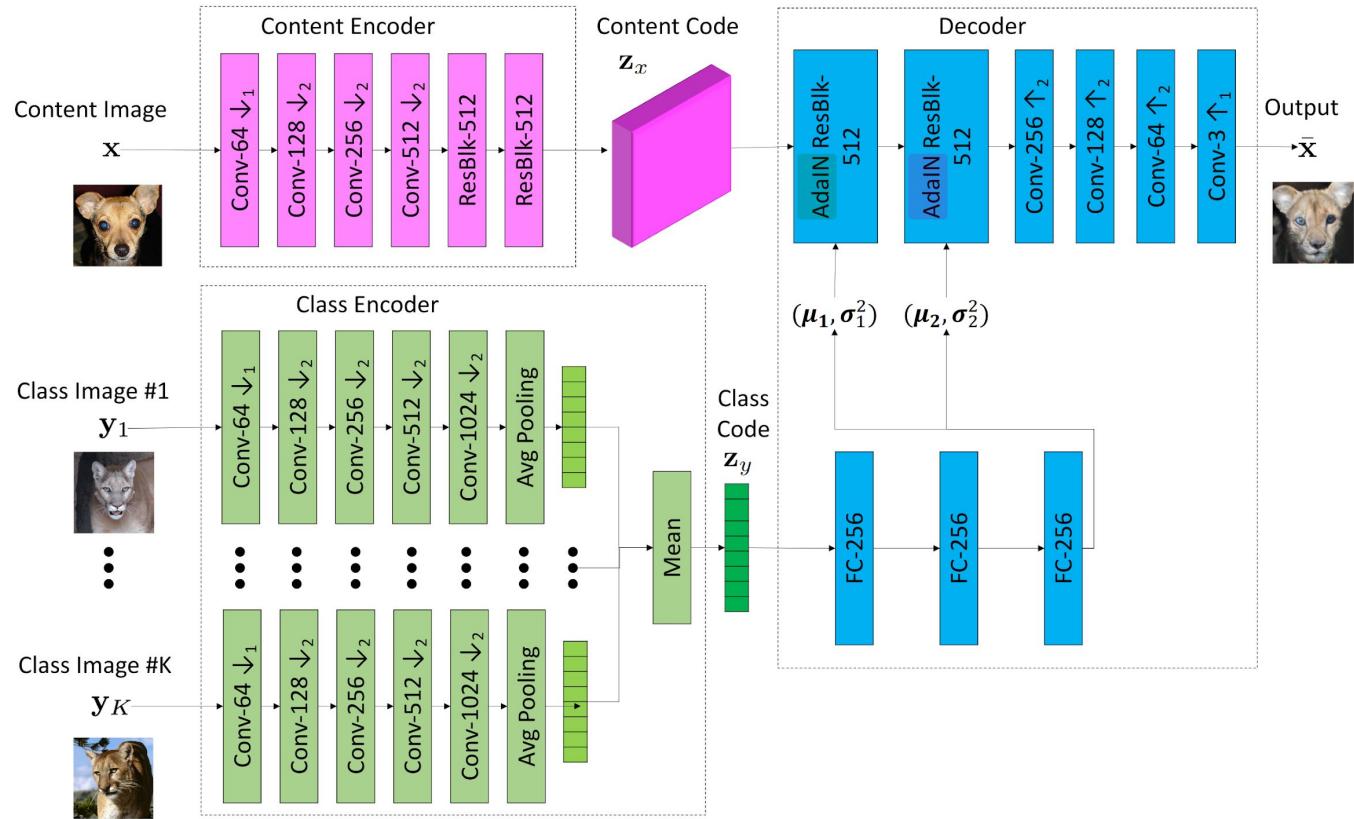




Input x	Class image y_1	Class image y_2	StarGAN -Unfair-5	StarGAN -Fair-5	FUNIT-5 \bar{x}
--------------	----------------------	----------------------	----------------------	--------------------	----------------------







Loss of FUNIT

$$\min_D \max_G \mathcal{L}_{\text{GAN}}(D, G) + \lambda_{\text{R}} \mathcal{L}_{\text{R}}(G) + \lambda_{\text{F}} \mathcal{L}_{\text{FM}}(G)$$

$$\begin{aligned}\mathcal{L}_{\text{GAN}}(G, D) = & E_{\mathbf{x}} [-\log D^{c_x}(\mathbf{x})] + \\ & E_{\mathbf{x}, \{\mathbf{y}_1, \dots, \mathbf{y}_K\}} [\log (1 - D^{c_y}(\bar{\mathbf{x}}))]\end{aligned}$$

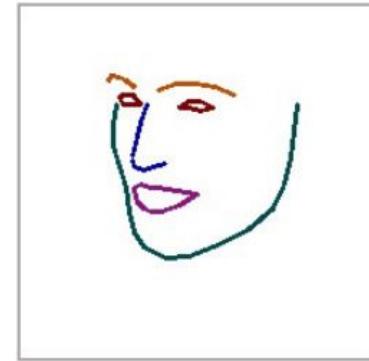
$$\mathcal{L}_{\text{R}}(G) = E_{\mathbf{x}} [||\mathbf{x} - G(\mathbf{x}, \{\mathbf{x}\})||_1^1]$$

$$\mathcal{L}_{\text{F}}(G) = E_{\mathbf{x}, \{\mathbf{y}_1, \dots, \mathbf{y}_K\}} [||D_f(\bar{\mathbf{x}})) - \sum_k \frac{D_f(\mathbf{y}_k)}{K}||_1^1].$$

Few-Shot Adversarial Learning of Realistic Neural Talking Head Models

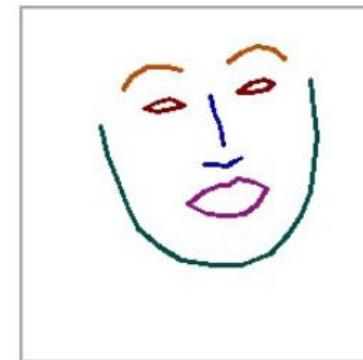
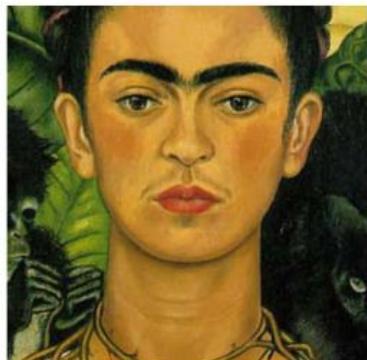
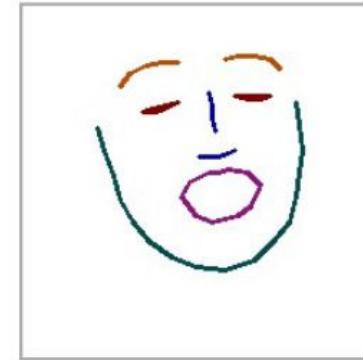
Egor Zakharov^{1,2} Aliaksandra Shysheya^{1,2} Egor Burkov^{1,2} Victor Lempitsky^{1,2}

¹Samsung AI Center, Moscow ²Skolkovo Institute of Science and Technology



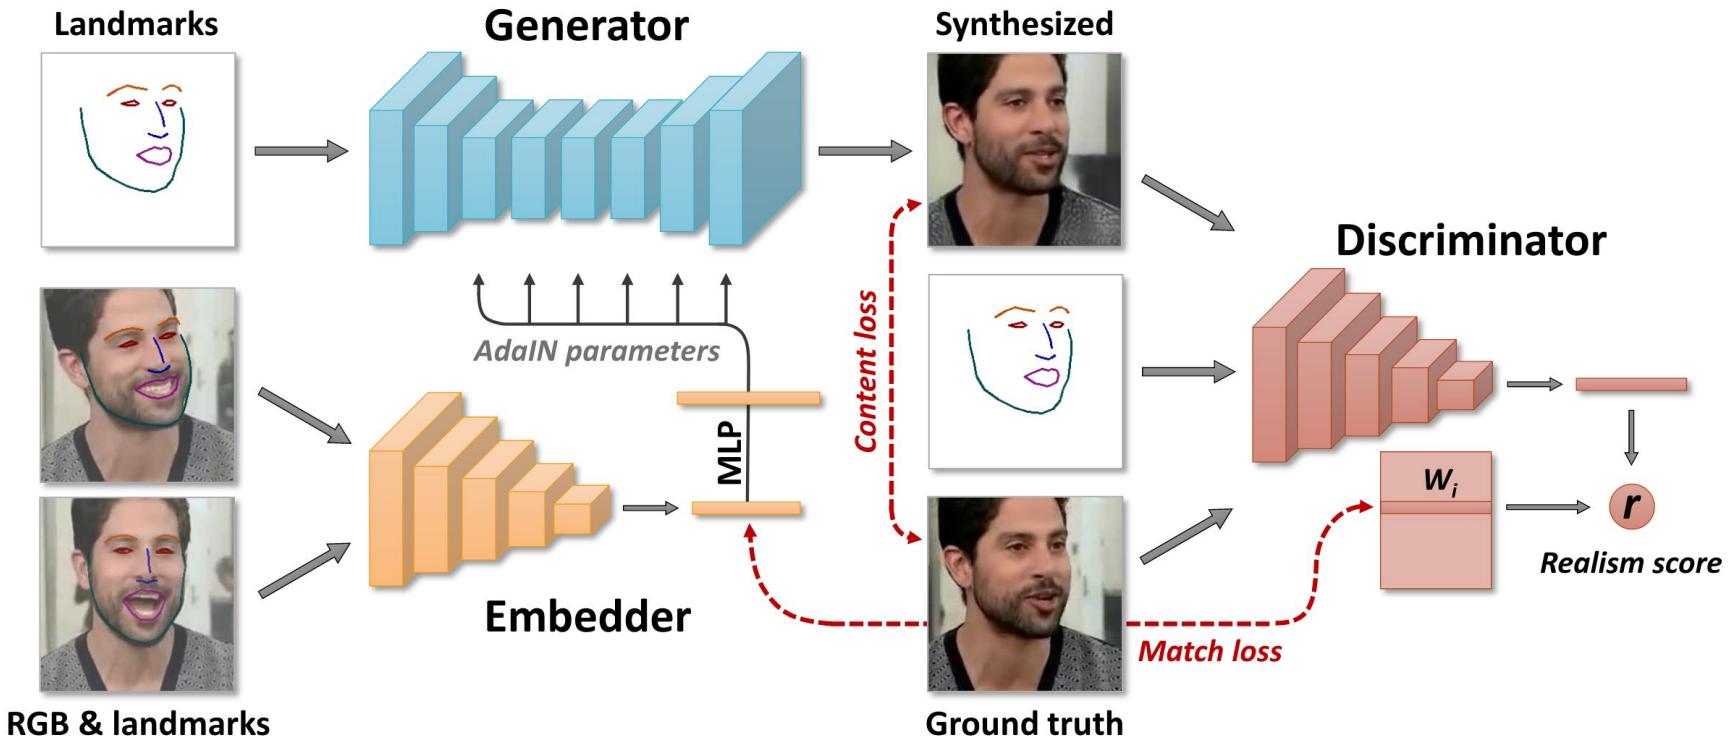
Source

Target → Landmarks → Result



Source

Target → Landmarks → Result



LOSS(G, E)

$$\begin{aligned}\mathcal{L}(\phi, \psi, \mathbf{P}, \theta, \mathbf{W}, \mathbf{w}_0, b) = & \mathcal{L}_{\text{CNT}}(\phi, \psi, \mathbf{P}) + \\ & \mathcal{L}_{\text{ADV}}(\phi, \psi, \mathbf{P}, \theta, \mathbf{W}, \mathbf{w}_0, b) + \mathcal{L}_{\text{MCH}}(\phi, \mathbf{W})\end{aligned}$$

$$\begin{aligned}\mathcal{L}_{\text{ADV}}(\phi, \psi, \mathbf{P}, \theta, \mathbf{W}, \mathbf{w}_0, b) = & \\ -D(\hat{\mathbf{x}}_i(t), \mathbf{y}_i(t), i; \theta, \mathbf{W}, \mathbf{w}_0, b) + \mathcal{L}_{\text{FM}}\end{aligned}$$

$$\mathcal{L}_{\text{FM}}(G, D_k) = \mathbb{E}_{(\mathbf{s}, \mathbf{x})} \sum_{i=1}^T \frac{1}{N_i} [||D_k^{(i)}(\mathbf{s}, \mathbf{x}) - D_k^{(i)}(\mathbf{s}, G(\mathbf{s}))||_1]$$

$$\begin{aligned}D(\hat{\mathbf{x}}_i(t), \mathbf{y}_i(t), i; \theta, \mathbf{W}, \mathbf{w}_0, b) = & \\ V(\hat{\mathbf{x}}_i(t), \mathbf{y}_i(t); \theta)^T (\mathbf{W}_i + \mathbf{w}_0) + b\end{aligned}$$

Loss Discriminator

$$\begin{aligned}\mathcal{L}_{\text{DSC}}(\phi, \psi, \mathbf{P}, \theta, \mathbf{W}, \mathbf{w}_0, b) = \\ \max(0, 1 + D(\hat{\mathbf{x}}_i(t), \mathbf{y}_i(t), i; \phi, \psi, \theta, \mathbf{W}, \mathbf{w}_0, b)) + \\ \max(0, 1 - D(\mathbf{x}_i(t), \mathbf{y}_i(t), i; \theta, \mathbf{W}, \mathbf{w}_0, b)).\end{aligned}$$

Summary

- Few shot learning 的一種方式是
 - 將資料編碼
 - 利用少量資料的編碼來設定模型
 - GAN 可以用這種方式來達到 Few shot learning



Data Augmentation GAN

Data Augmentation GAN

Antreas Antoniou, Amos Storkey, Harrison Edwards,
“Data Augmentation Generative Adversarial
Networks”



Motamed, Saman, and Farzad Khalvati. “Inception Augmentation Generative Adversarial Network.” arXiv preprint arXiv:2006.03622

DAGAN

作用在 VGG-Face 的結果(Test Accuracy)如下

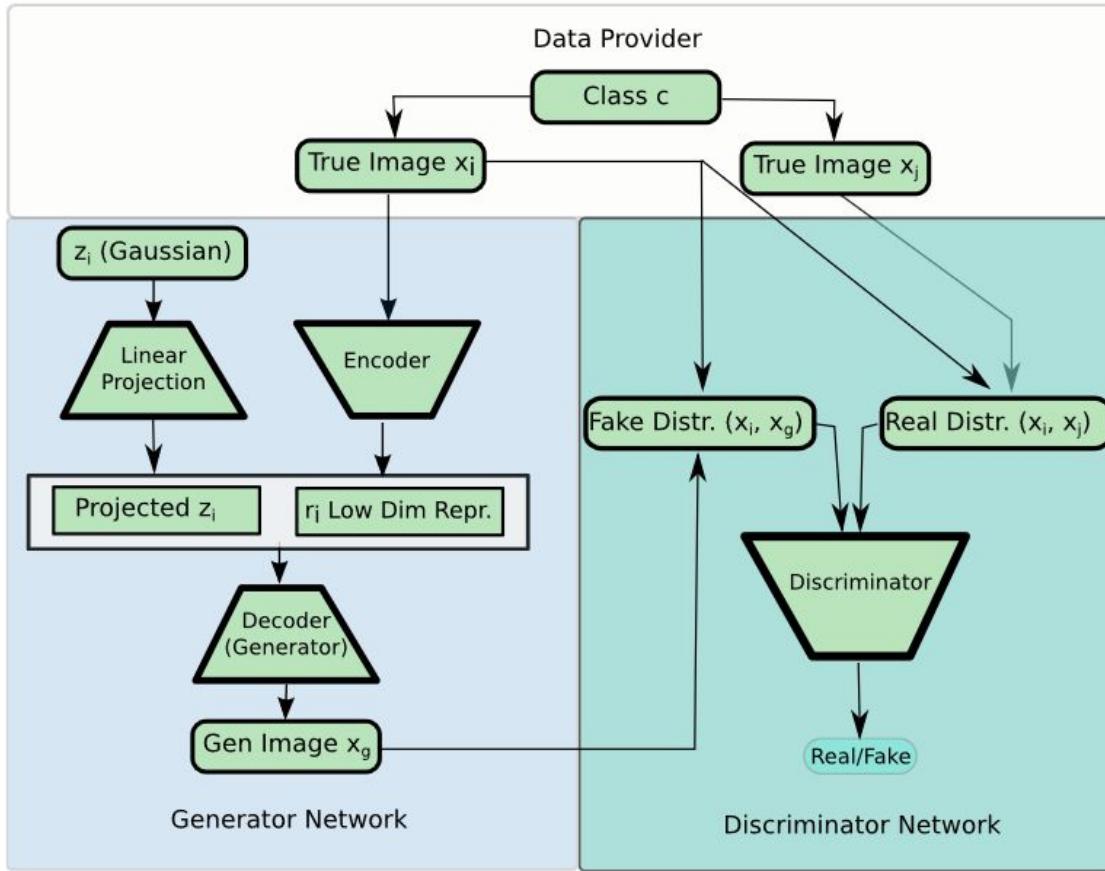
# Sample...	# Standard	# DAGAN
5	4.4%	12.6%
15	39.3%	42.9%
25	58%	58.5%

直接用 GAN 來 augmentation

利用 GAN 來學 augmentation.

- 紿定輸入圖片
- 先取得 representation ,
- 然後利用生成模型 和 noise 來生出

這裡的生成模型可以視為 unsupervised meta-learning.











DAGAN

作用在 VGG-Face 的結果(Test Accuracy)如下

# Sample...	# Standard	# DAGAN
5	4.4%	12.6%
15	39.3%	42.9%
25	58%	58.5%

Inception-Augmentation GAN

架構類似 DAGAN, 也是由輸入的圖片和 noise 產生新圖片

- 但包含較為先進的結構, 像是attention, inception blocks

輸入的圖片先

- 經過 CNN 和 attention 降低維度
- 和 noise concate
- 然後再進入後面的網路。

訓練 GAN 的時候:

- 輸入 x , 然後希望生出來的圖片盡量和 x 不同(以 SSIM 衡量)
- 但又符合原來的 distribution (conditional GAN 的意義下)。

當 GAN 訓練完後

- 利用 AnoGAN 的方式來判斷異常。

原始資料有 3265 張的肺部圖片。

- 在沒有 augmentation 的情形下, AUC=0.83
- 使用 IAGAN, AUC=0.88.

這個方法主要的優點是，3265 張的數量不多

- 在其他 GAN 失效的情形下，利用圖片輔助生成，可以讓GAN 更容易生成。

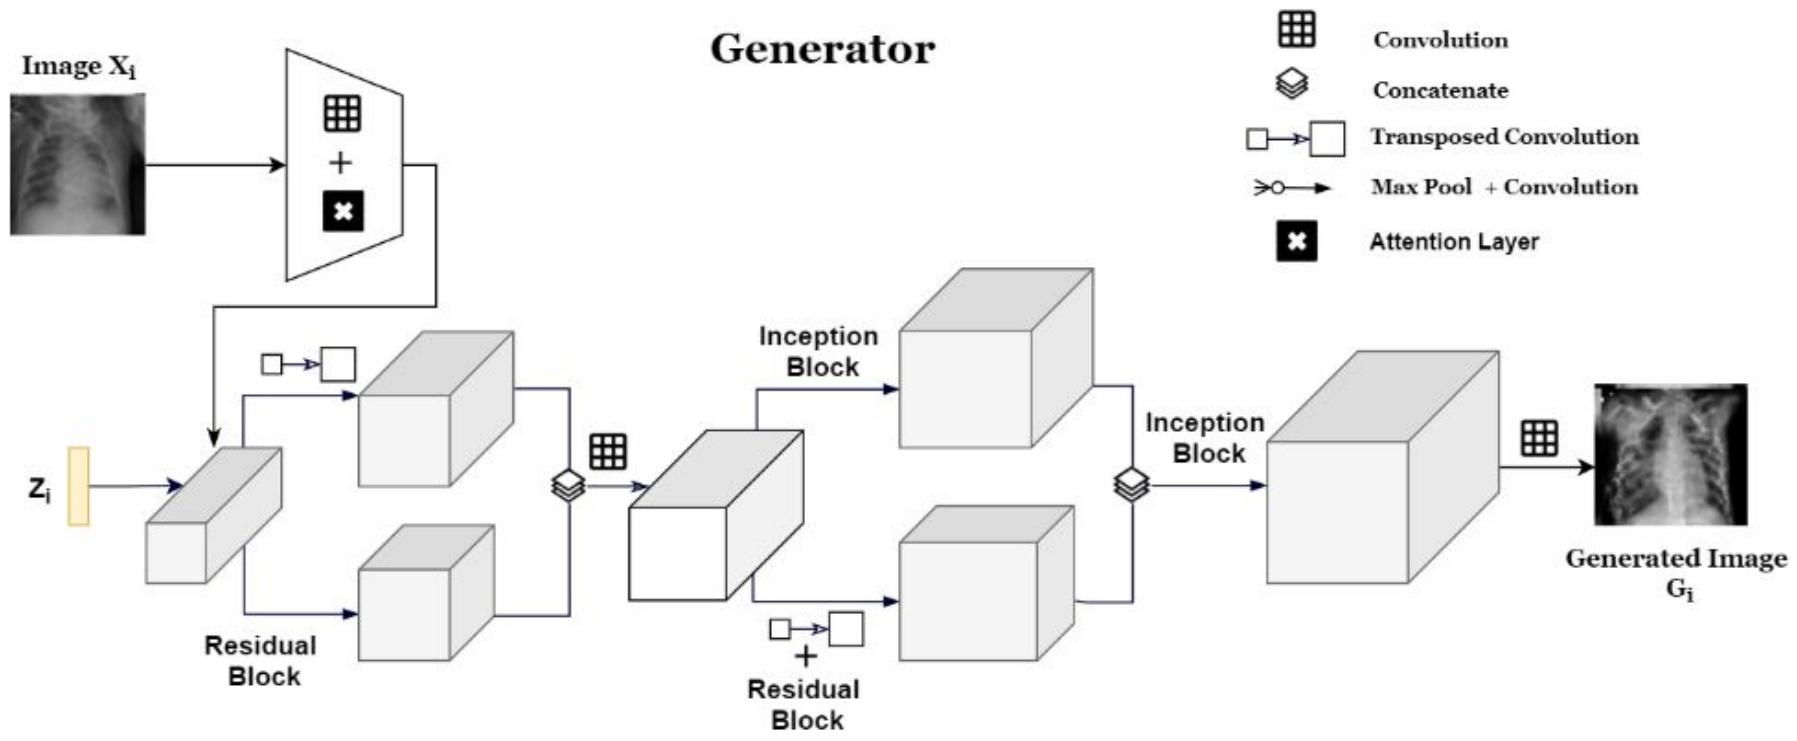


Figure 1: IAGAN's Generator Architecture

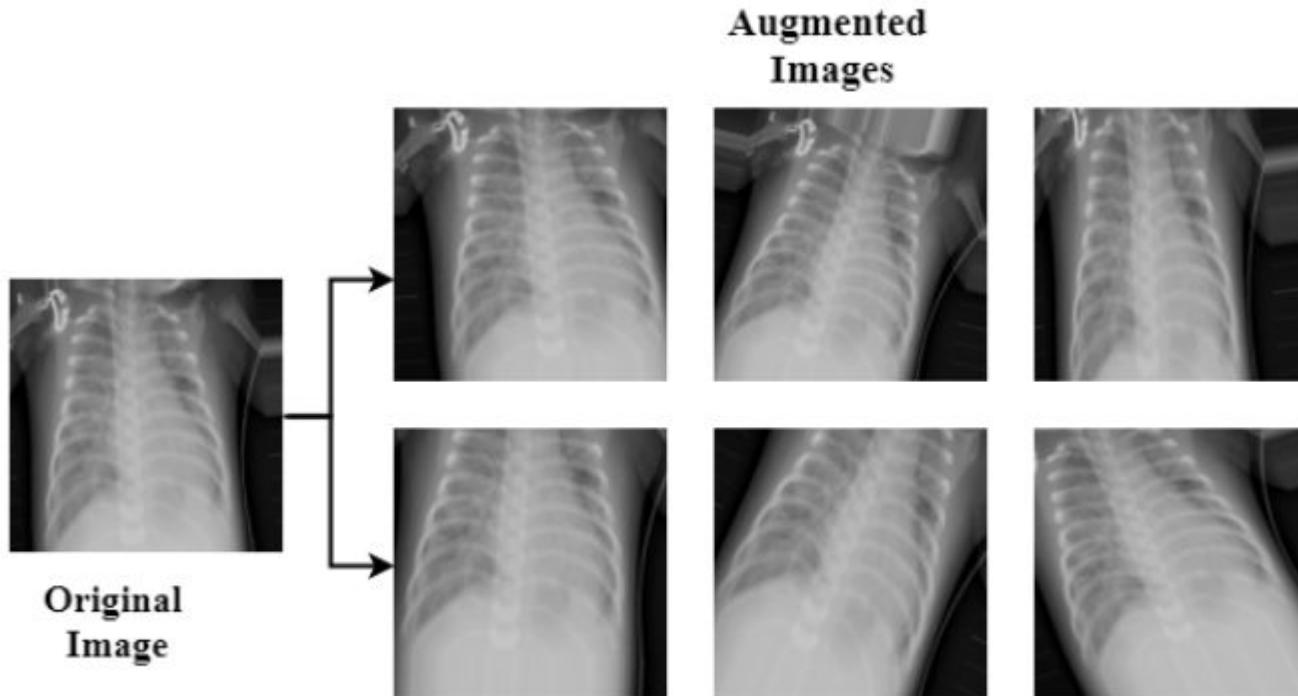
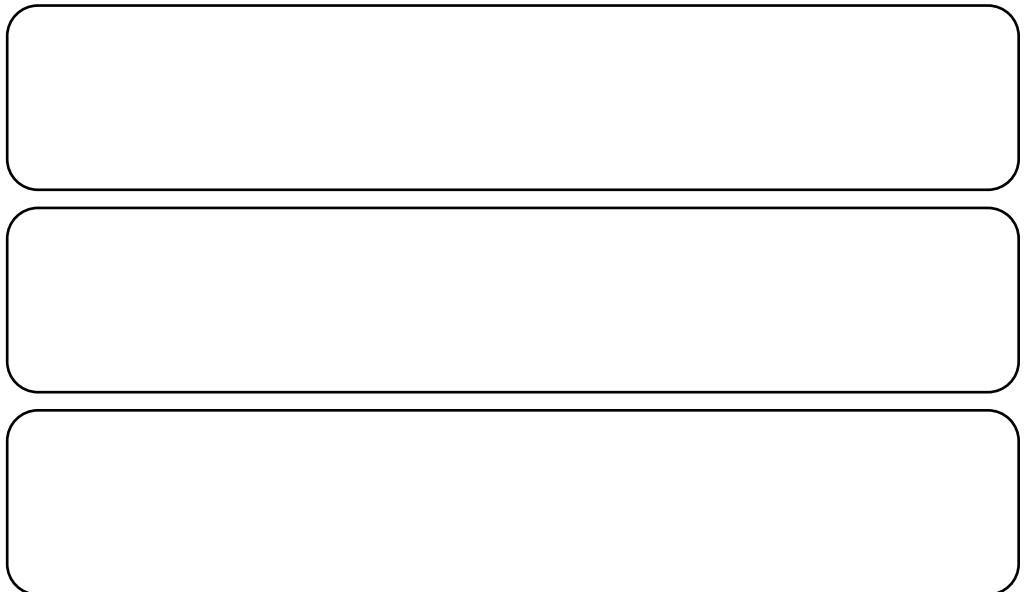
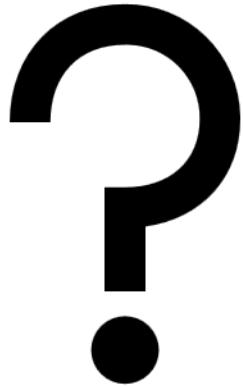


Figure 6: Traditional augmentation output sample

Summary

- 更多不同的 GAN 使用方式
- 利用圖到圖來
 - 降低訓練集數量的需求
 - 提升生成圖片的品質

StyleGAN2-ADA 特點是





Review (semi-supervised GAN)

ACGAN

- <https://arxiv.org/abs/1610.09585>

InfoGAN

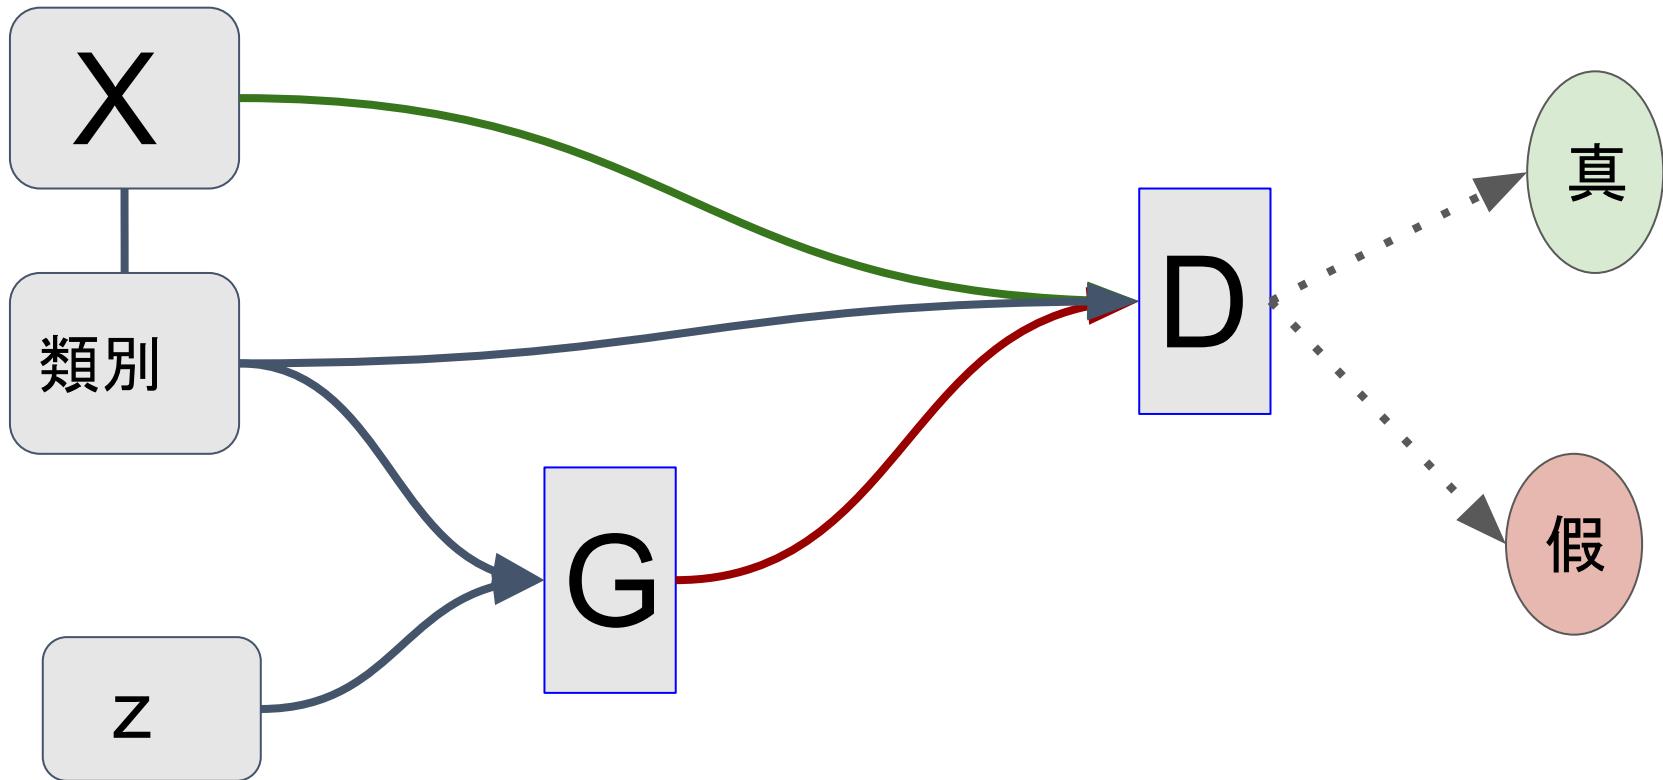
- <https://arxiv.org/abs/1606.03657>

Conditional GAN

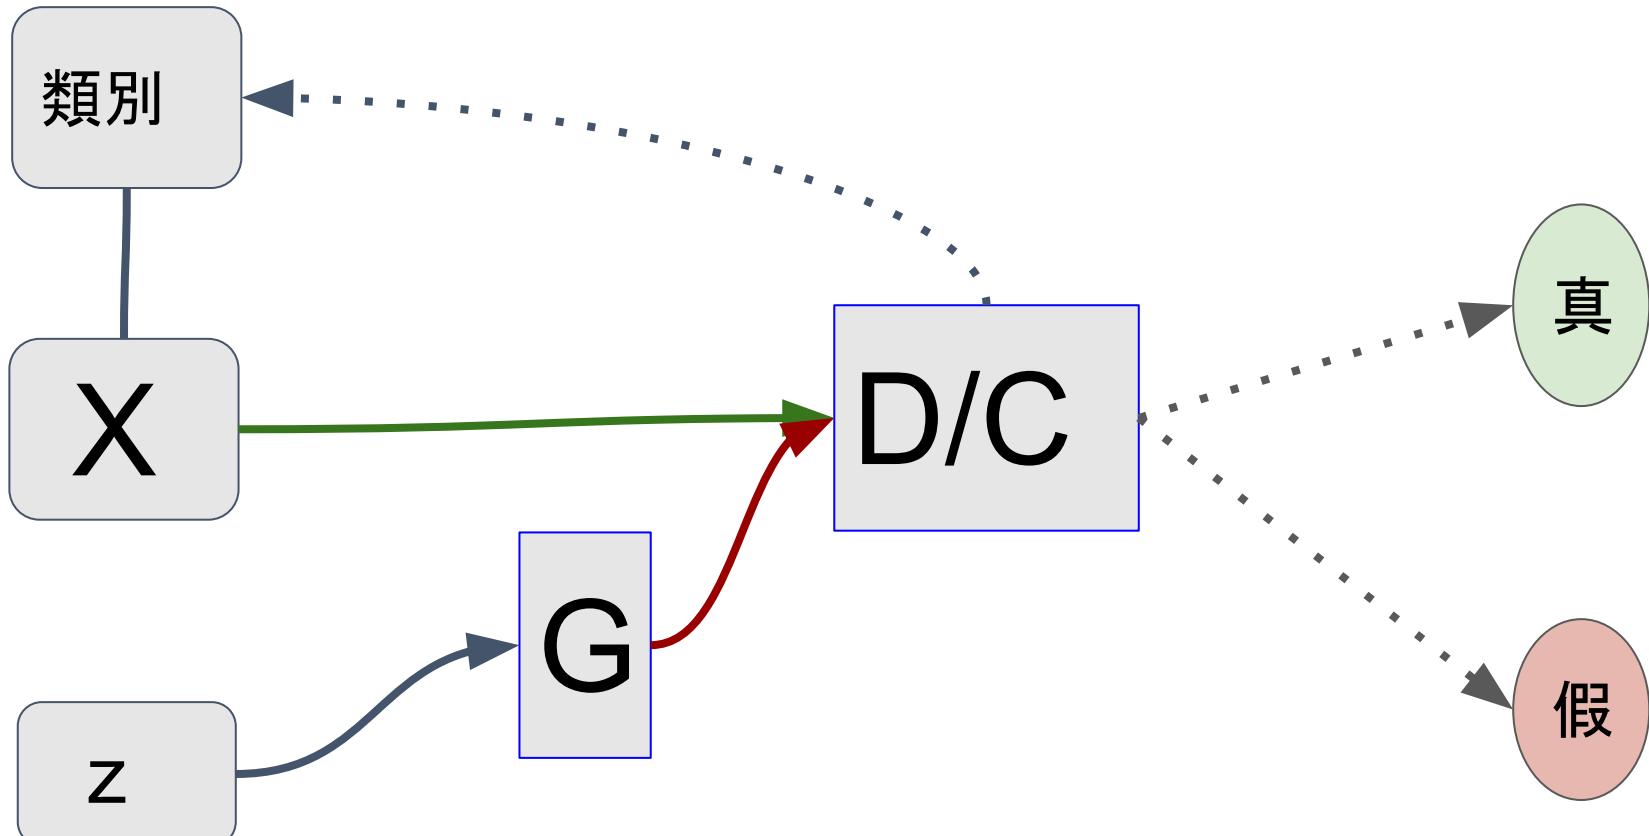
- <https://arxiv.org/abs/1411.1784>

Semi-supervised GAN

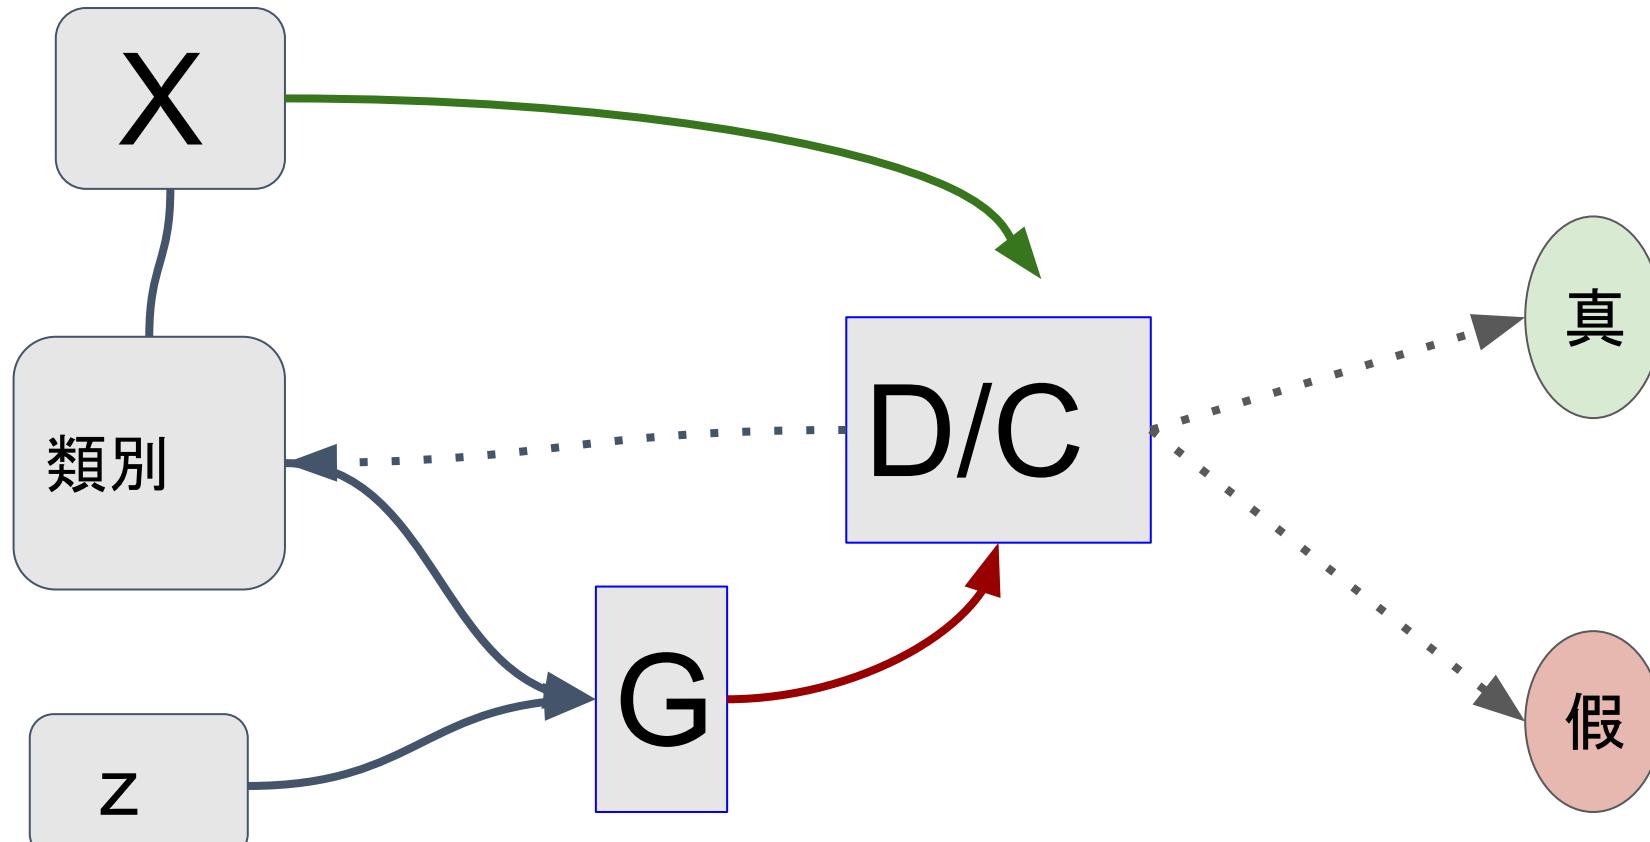
- <https://arxiv.org/abs/1606.01583>

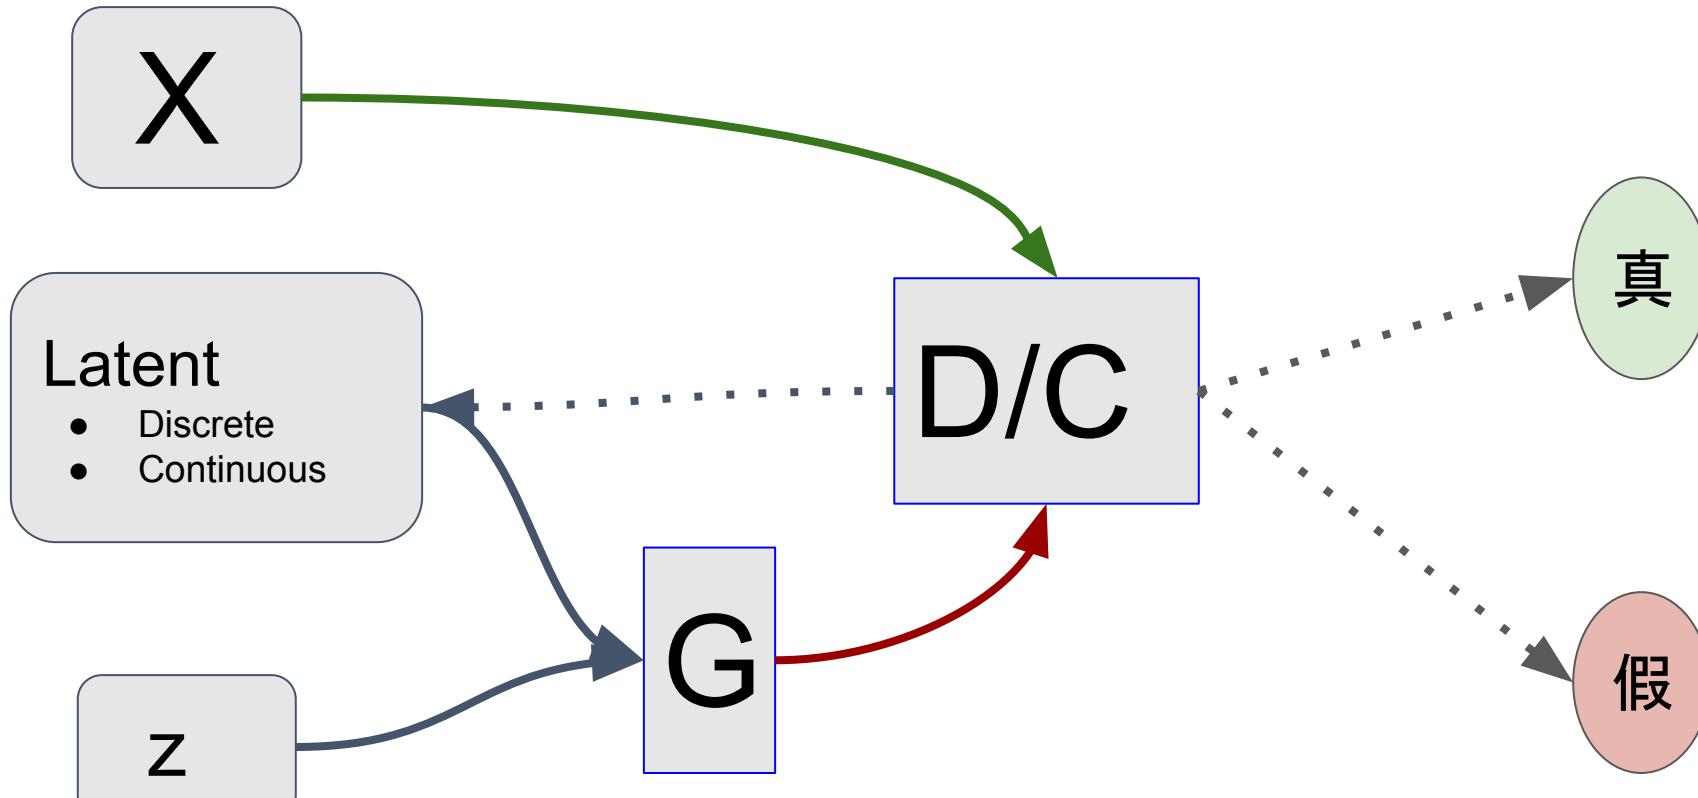


Conditional GAN



Semi-Supervised GAN





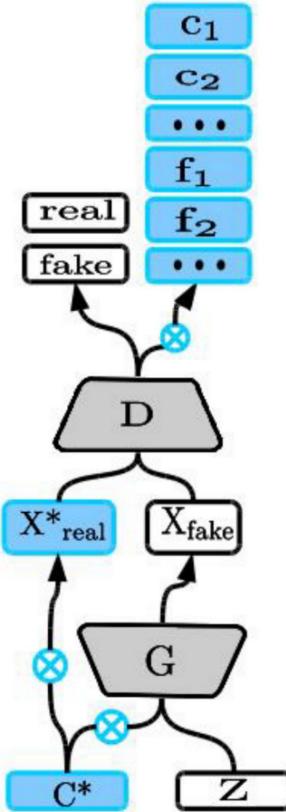
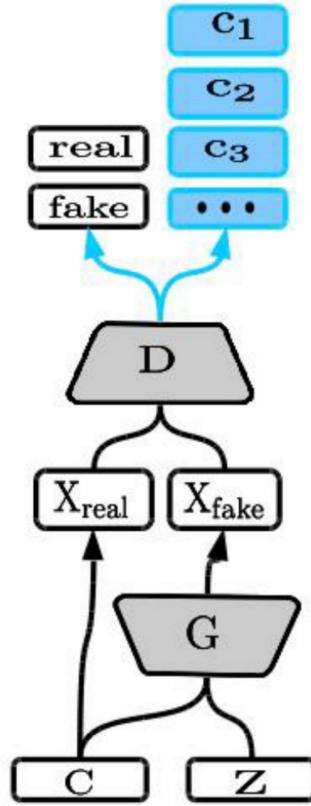
InfoGAN

GAN Based Few-shot Learning



Few-shot classifier GAN

- <http://www.animlife.com/publications/ijcnn18.pdf>

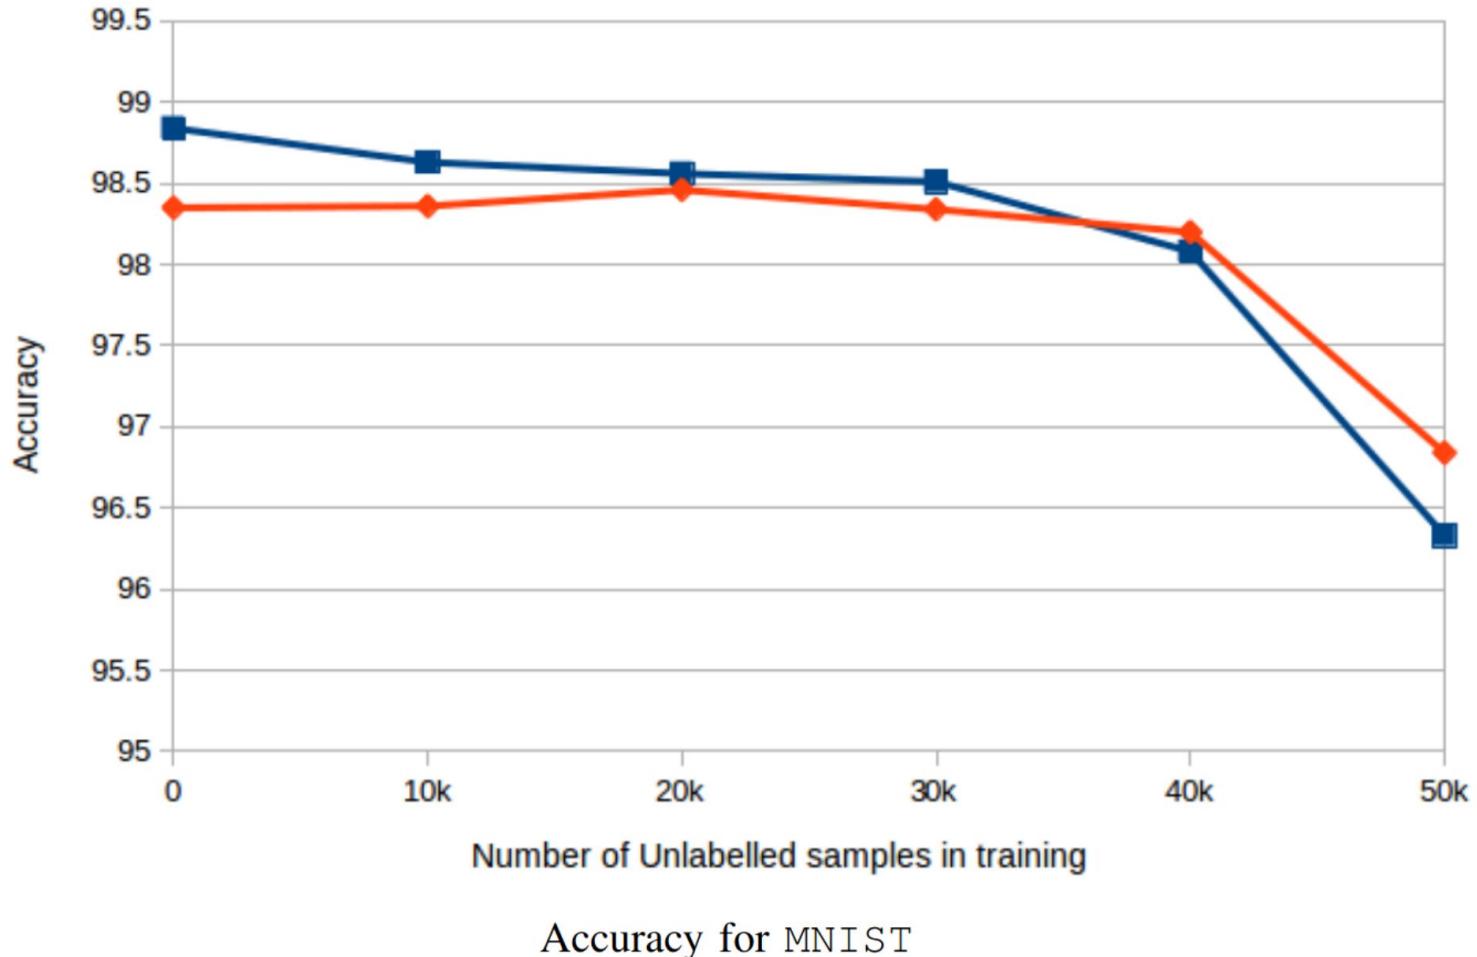


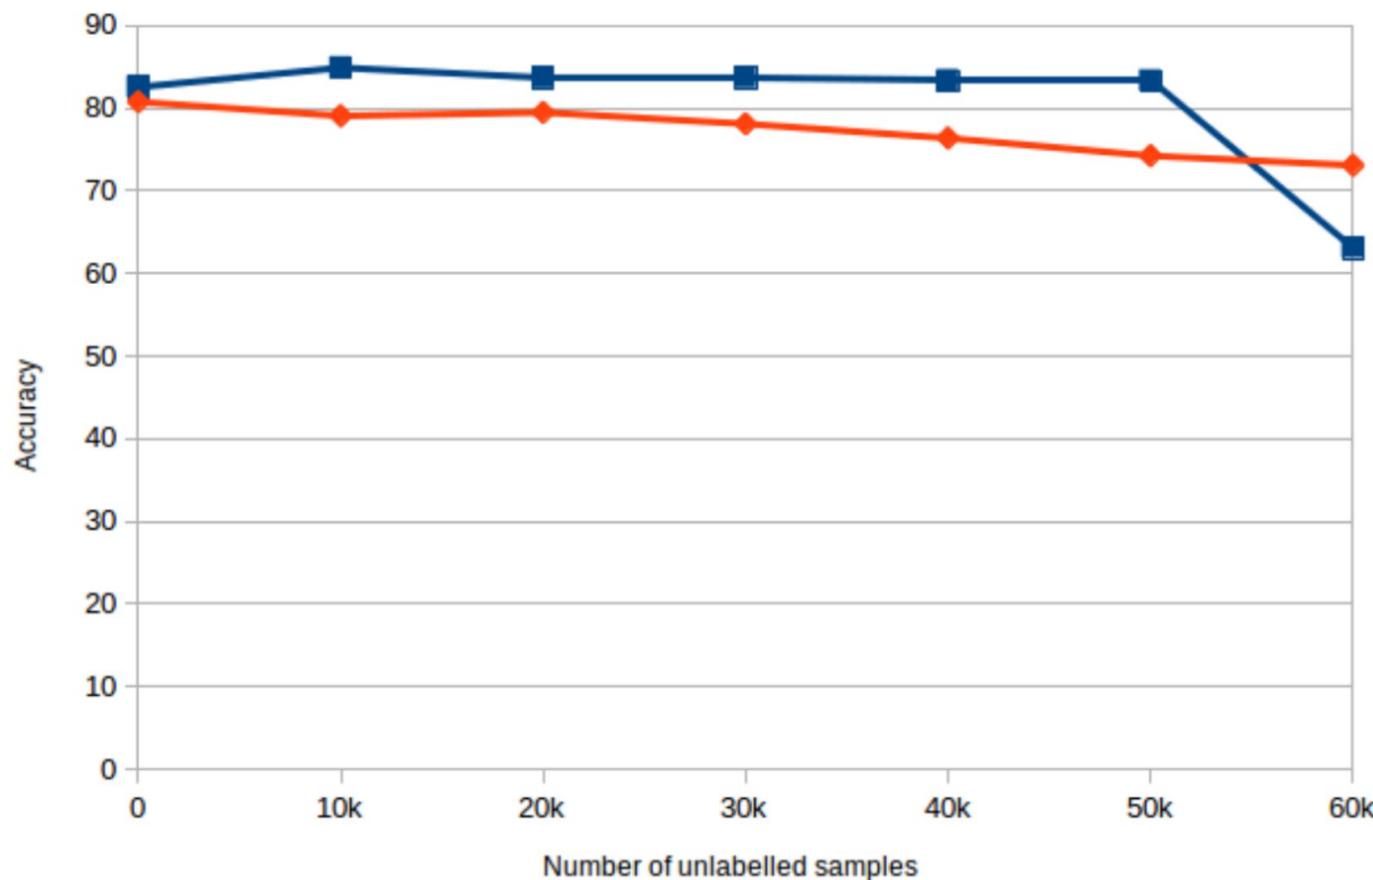
- 類似 ACGAN

- 但有時沒有 condition
- 利用網路內的 Switcher 選擇

Algorithm 1 Training Algorithm

```
1: procedure TRAIN(data_batches, label_batches)
2:   for e epochs do
3:     for k steps do
4:       Fetch next labeled mini batches
5:       Perform Stochastic Gradient Descent on D
6:       Perform Stochastic Gradient Descent on G
7:       Evaluate( $\mathcal{L}_s, \mathcal{L}_c$ )
8:       Update D and G losses
9:     end for
10:    for j steps do
11:      Fetch next unlabeled data mini batches
12:      Perform Stochastic Gradient Descent on D
13:      Perform Stochastic Gradient Descent on G
14:      Evaluate( $V_s$ )
15:      Update D and G losses
16:    end for
17:  end for
18: end procedure
```

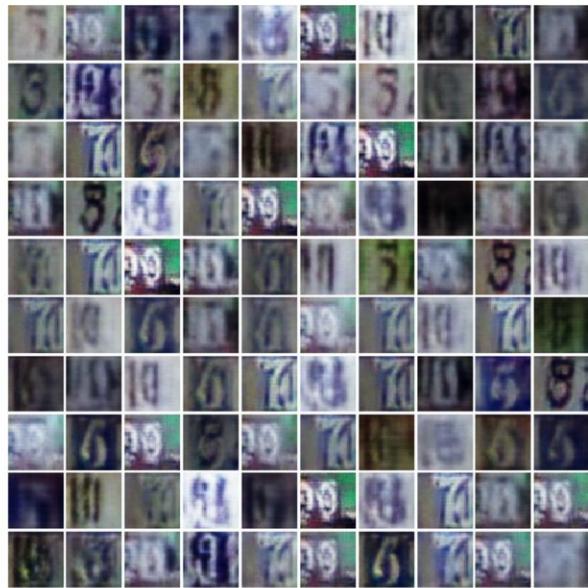




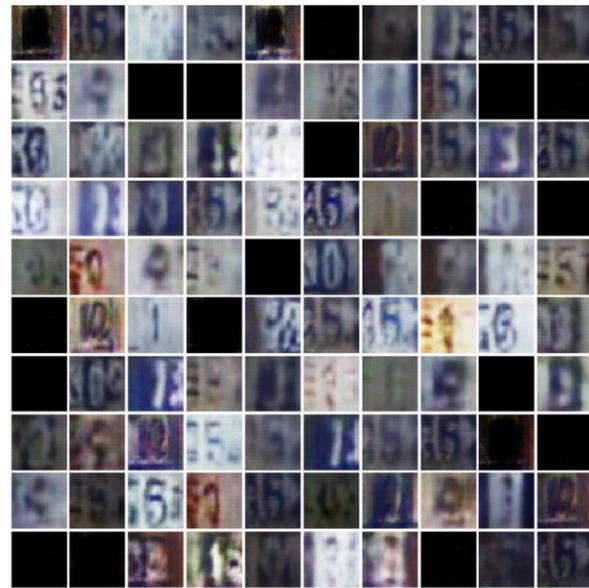
Accuracy for SVHN



Original sample



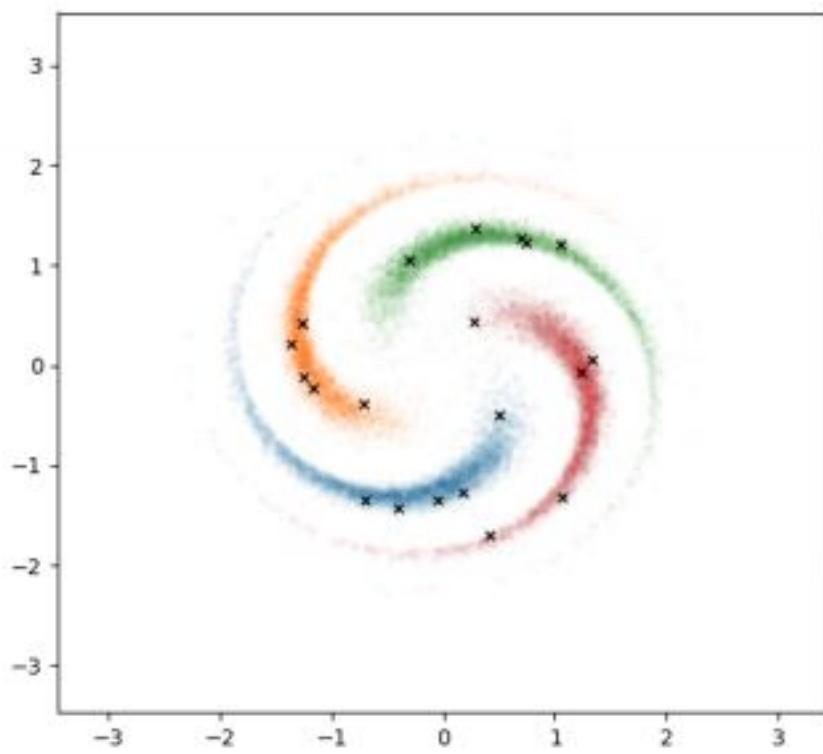
Generated sample (single fake class)

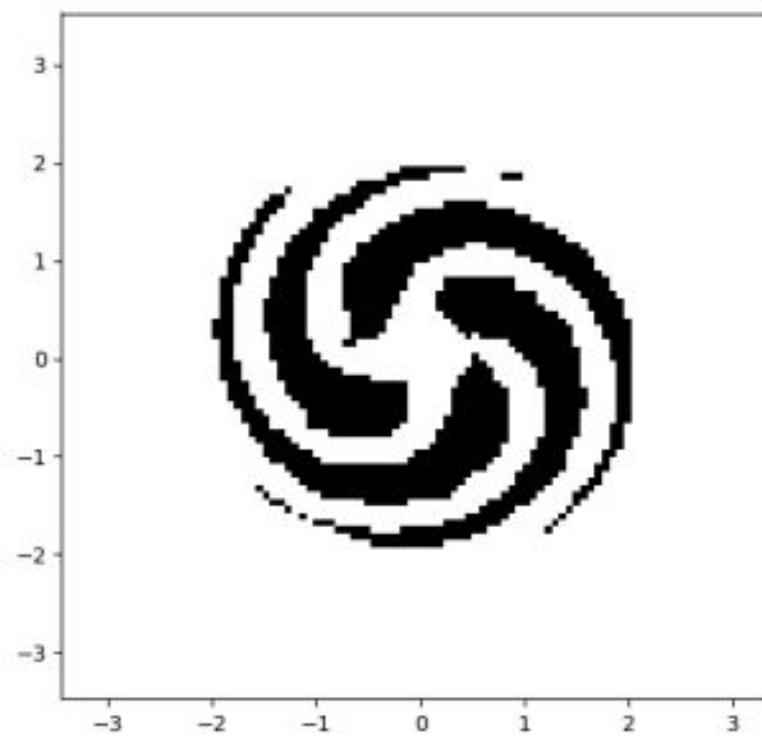
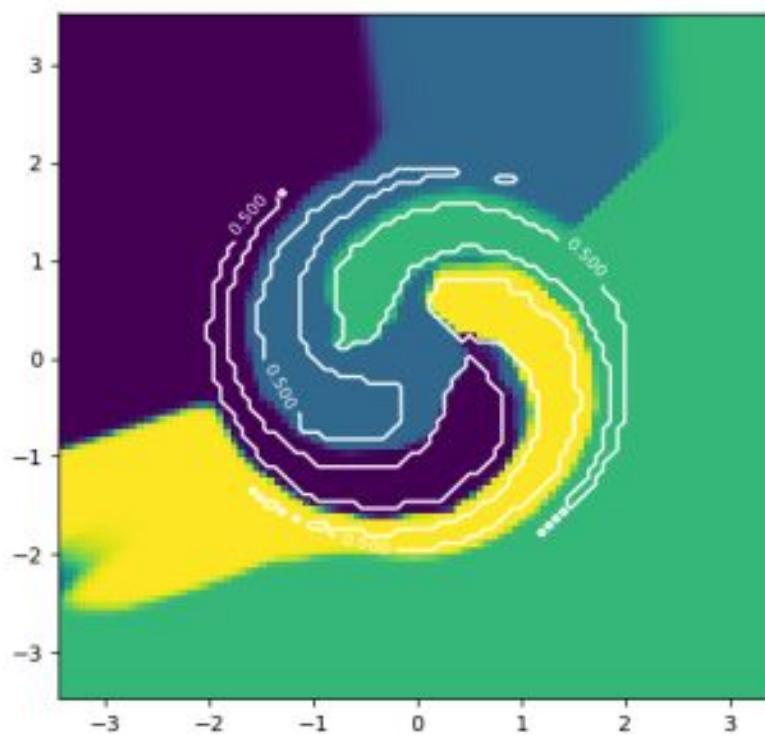


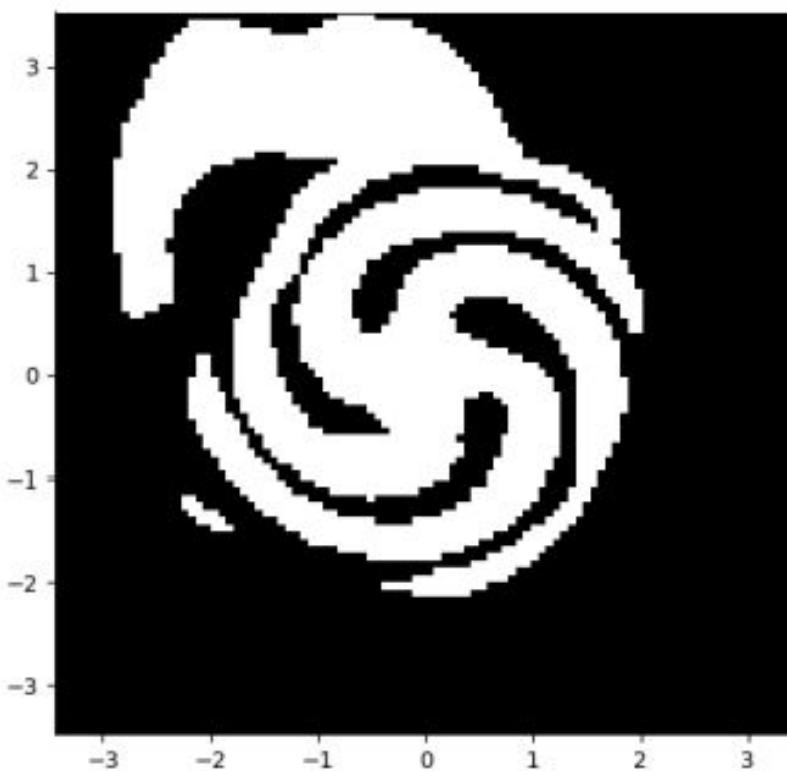
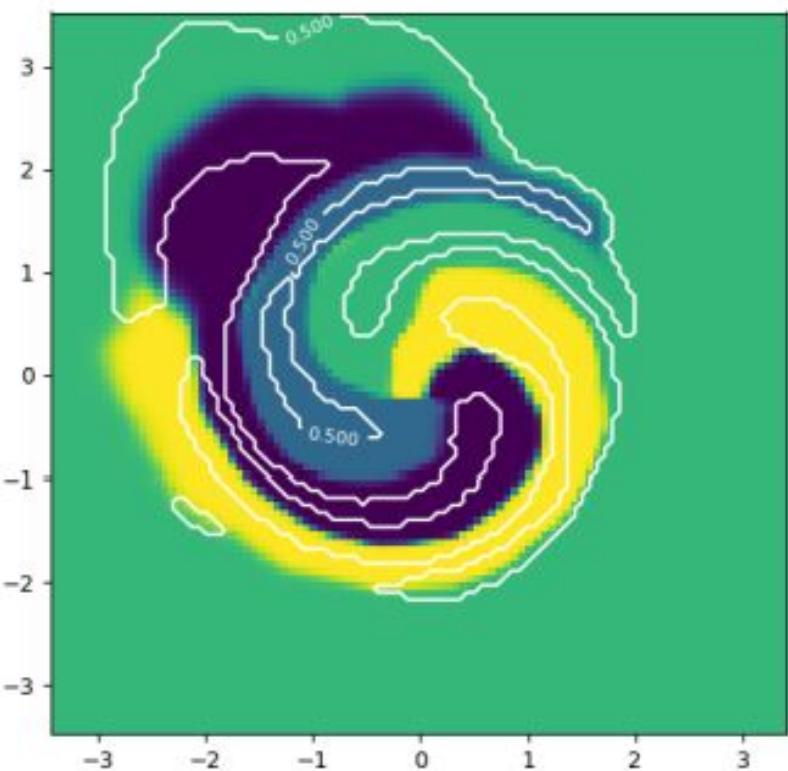
Generated sample (multiple fake classes)

Good semi-supervised learning that requires a bad GAN.

- <https://proceedings.neurips.cc/paper/2017/file/79514e888b8f2acac68738d0cbb803e-Paper.pdf>







Summary

- GAN 可以靈活應用未標記的資料
- 即使 GAN 生成的品質不好，也可能有幫助
- GAN 生出的東西當成是瑕疵可能可以在真實資料附近建立更緊密的邊界

A woman with short, multi-colored hair (pink, yellow, blue) and visible tattoos on her neck and chest. She is wearing a light-colored, mesh-style t-shirt. The background is dark with a grid of glowing blue binary digits (0s and 1s).

GAN powerd Meta-Learning

MetaGAN

- <https://papers.nips.cc/paper/2018/hash/4e4e53aa080247bc31d0eb4e7aeb07a0-Abstract.html>

Settings of Few Shot Learning

- K -shot N -way supervised learning
 - Sample-level Semi-supervised Learning
 - Task-level Semi-supervised Learning

Task

- Predict y given x
- Support set $S_T = (S_T^s, S_T^u)$
 - $S_T^s = \{(x_i, y_i)\}$: labeled data
 - S_T^u : unlabeled data
- Query set $Q_T = (Q_T^s, Q_T^u)$

Dataset

- Training: contains tasks
 - (Support, Query)
- Testing: contains tasks
 - (Support, Query)

Sample-level Semi-supervised Learning

- Allow some training samples to be unlabeled [within a task](#)
- These training samples can either come from the same classes as the labeled samples
- or come from different "distractor" classes

Task-level Semi-supervised Learning

- also allow purely unsupervised task
 - in which both support and query samples are all unlabeled
- can be very natural in practice

Discriminator

- Use **decent few-shot learners**, e.g.
 - MAML
 - Relation Network
- Add an additional output, as in semi-supervised learning
 - So there are $N+1$ outputs

LOSS

$$\mathcal{L}_D^{\mathcal{T}} = \mathcal{L}_{\text{supervised}} + \mathcal{L}_{\text{unsupervised}},$$

$$\mathcal{L}_{\text{supervised}} = \mathbb{E}_{\mathbf{x}, y \sim Q_{\mathcal{T}}^s} \log p_D(y|\mathbf{x}, y \leq N)$$

$$\mathcal{L}_{\text{unsupervised}} = \mathbb{E}_{\mathbf{x} \sim Q_{\mathcal{T}}^u} \log p_D(y \leq N|\mathbf{x}) + \mathbb{E}_{\mathbf{x} \sim p_G^{\mathcal{T}}} \log p_D(N+1|\mathbf{x})$$

Generator

- Minimize

$$L_G^{\mathcal{T}}(D) = -\mathbb{E}_{\mathbf{x} \sim p_G^{\mathcal{T}}} [\log(p_D(y \leq N | \mathbf{x}))].$$

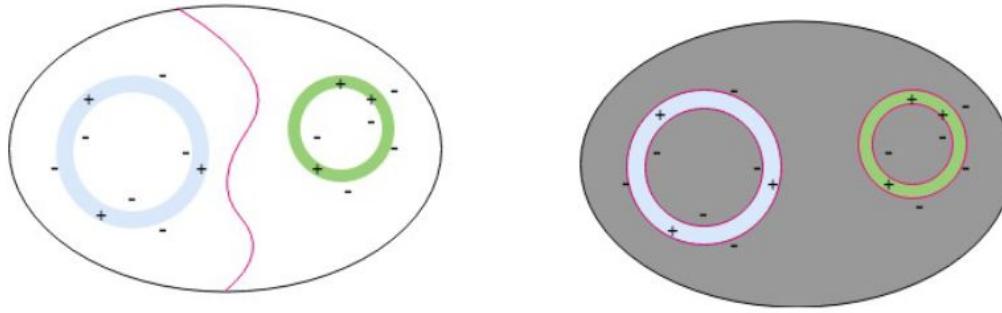
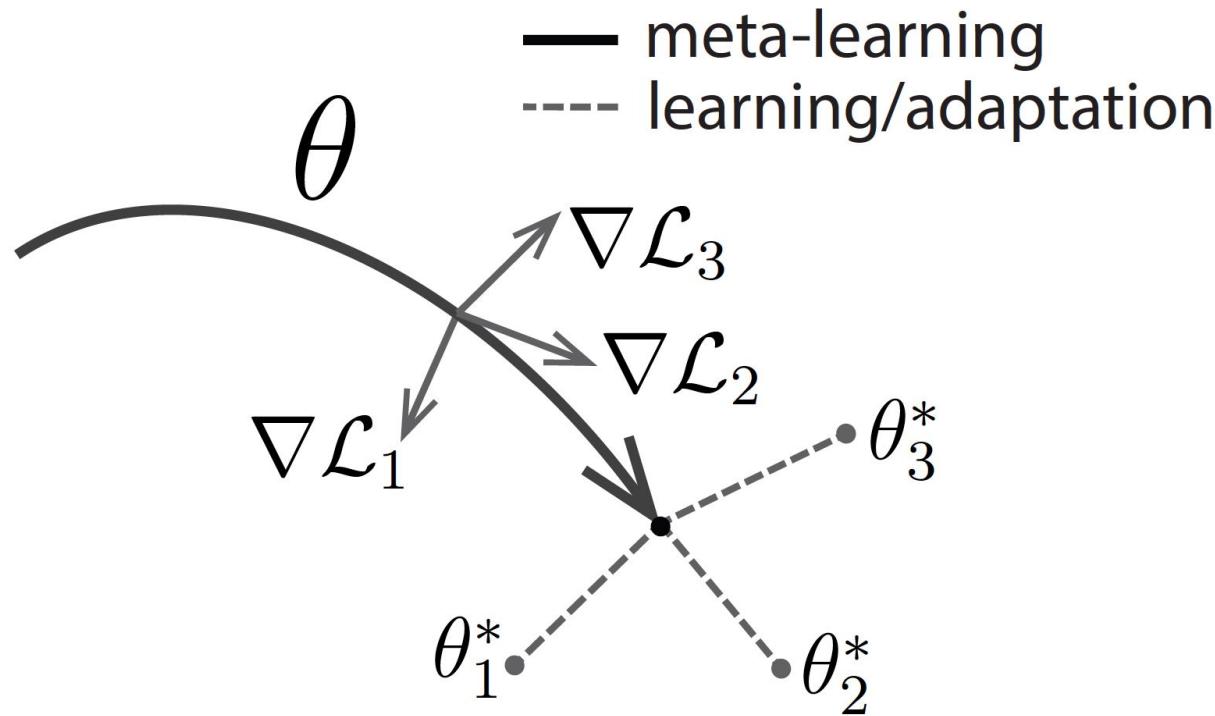


Figure 1: Left: decision boundary without metaGAN. Right: decision boundary with metaGAN. We use red curves to denote the decision boundary. Blue area in figure represents class A, green area represents class B, and gray area represents fake class. We use $+$ to denote real samples and $-$ to denote fake samples generated.

MAML



GAN With MAML

- Use the MAML model as **discriminator**
- For a specific task, update the model
- Update the discriminator according to the loss

$$\ell_D^T = -\mathbb{E}_{\mathbf{x}, y \sim S_T^s} \log p_D(y|\mathbf{x}, y \leq N) - \mathbb{E}_{\mathbf{x} \sim S_T^u} \log p_D(y \leq N|\mathbf{x}) - \mathbb{E}_{\mathbf{x} \sim p_G^T} \log p_D(N+1|\mathbf{x})$$

- Train the generator using the **adapted** discriminator

Algorithm 1 MetaGAN with MAML

$G(\mathbf{z}, \mathcal{T})$: Generator network parameterized by θ_g .
 $D(x)$: Discriminator network. parameterized by θ_d .

Initialize θ_g, θ_d randomly.

while not done **do**

 Sample a batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$. ▷ Discriminator Update

for all \mathcal{T}_i **do**

 Get K real samples $\mathcal{D}_r = \{\mathbf{x}^{(i)}, y^{(i)}\}$ from \mathcal{T}_i .

 Sample K generated samples $\mathcal{D}_f = \{\mathbf{x}^{(j)}\} = G(\mathbf{z}^{(j)}, \mathcal{T}_i)$ from $G(\mathbf{z}, \mathcal{T}_i)$.

 Evaluate discriminator loss $\ell_D^{\mathcal{T}_i}$ with D_r and D_f .

 Compute adapted discriminator parameters $\theta'_{d_i} = \theta_d - \alpha \nabla_{\theta_d} \ell_D^{\mathcal{T}_i}$.

end for

 Update θ_d using loss \mathcal{L}_D

 Sample a batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$. ▷ Generator Update

for all \mathcal{T}_i **do**

 Sample K generated samples $\mathcal{D}_f = \{\mathbf{x}^{(j)} = G(\mathbf{z}^{(j)}, \mathcal{T}_i)\}$ from $G(\mathbf{z}, \mathcal{T}_i)$.

 Compute adapted discriminator parameters $\theta'_{d_i} = \theta_d - \alpha \nabla_{\theta_d} L_D$.

 Compute generator loss gradient $\nabla_{\theta_g} L_G^{\mathcal{T}_i}$ with the adapted discriminator.

end for

 Update generator parameters θ_g with accumulated generator loss gradients.

end while

Relation Network

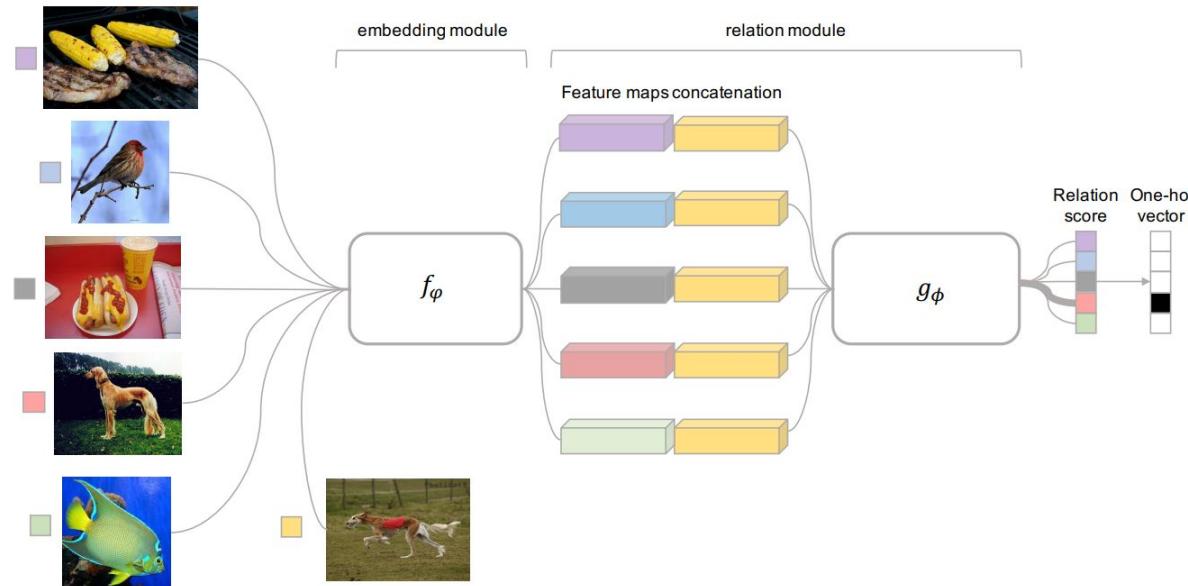


Figure 1: Relation Network architecture for a 5-way 1-shot problem with one query example.

GAN With Relation Network

- Use softmax classification on relevance score between query set and support set

$$p_D(y = k | \mathbf{x}_j) = \frac{\exp(r_{k,j})}{1 + \sum_{i=1}^N \exp(r_{i,j})}$$

Model	5-way Acc.		20-way Acc.	
	1-shot	5-shot	1-shot	5-shot
Neural Statistician	98.1	99.5	93.2	98.1
Prototypical Nets	98.8	99.7	96.0	98.9
MAML	98.7 ± 0.4	99.9 ± 0.1	95.8 ± 0.3	98.9 ± 0.2
Ours: MetaGAN + MAML	99.1 ± 0.3	99.7 ± 0.21	96.4 ± 0.27	98.9 ± 0.18
Relation Net	99.6 ± 0.2	99.8 ± 0.1	97.6 ± 0.2	99.1 ± 0.1
Ours: MetaGAN + RN	99.67 ± 0.18	99.86 ± 0.11	97.64 ± 0.17	99.21 ± 0.1

Table 1: Few-shot classification results on Omniglot.

Model	5-way Acc.	
	1-shot	5-shot
Prototypical Nets	49.42 ± 0.78	68.20 ± 0.66
MAML(5 gradient steps)	48.70 ± 1.84	63.11 ± 0.92
MAML(5 gradient steps, first order)	48.07 ± 1.75	63.15 ± 0.91
MAML(1 gradient step, first order)	43.64 ± 1.91	58.72 ± 1.20
Ours: MetaGAN + MAML(1 step, first order)	46.13 ± 1.78	60.71 ± 0.89
Relation Net	50.44 ± 0.82	65.32 ± 0.7
Ours: MetaGAN + RN	52.71 ± 0.64	68.63 ± 0.67

Table 2: Few-shot classification results on Mini-Imagenet.

Model	Omniglot	Mini-Imagenet	
	1-shot 5-way	1-shot 5-way	5-shot 5-way
Prototypical Nets(Supervised)	94.62 ± 0.09	43.61 ± 0.27	59.08 ± 0.22
Semi-Supervised Inference(PN)	97.45 ± 0.05	48.98 ± 0.34	63.77 ± 0.20
Soft k-Means	97.25 ± 0.10	50.09 ± 0.45	64.59 ± 0.28
Soft k-Means+Cluster	97.68 ± 0.07	49.03 ± 0.24	63.08 ± 0.18
Masked Soft k-Means	97.52 ± 0.07	50.41 ± 0.31	64.39 ± 0.24
Ours: Relation Nets(Supervised)	94.81 ± 0.08	44.24 ± 0.24	58.72 ± 0.31
Ours: MetaGAN + RN	97.58 ± 0.07	50.35 ± 0.23	64.43 ± 0.27

Table 3: Sample-level Semi-Supervised Few-shot classification results on Omniglot and Mini-Imagenet.

Model	Omniglot	Mini-Imagenet
	1-shot 5-way	1-shot 5-way
Prototypical Net(Supervised)	93.66 ± 0.09	42.28 ± 0.32
Relation Net(Supervised)	93.82 ± 0.07	43.87 ± 0.20
Ours: MetaGAN + RN	97.12 ± 0.08	47.43 ± 0.27

Table 4: Task-level Semi-Supervised 1-shot classification results on Omniglot and Mini-Imagenet.

Summary

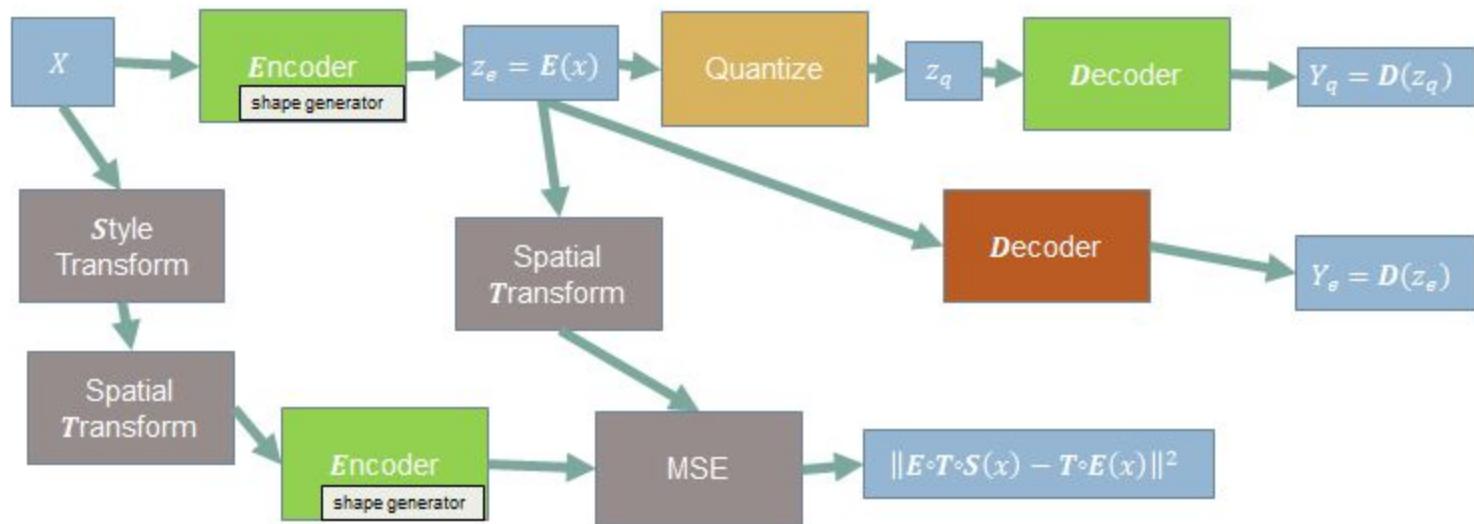
- GAN 可以搭配很多不同的架構
- 包含 Meta-learning
- GAN 在 Few-shot 中可以模仿未知的其他類別

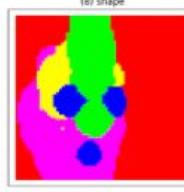
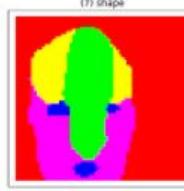
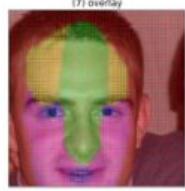
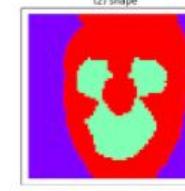
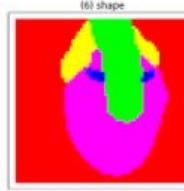
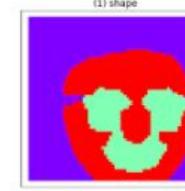
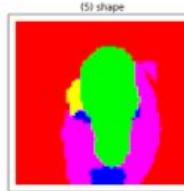
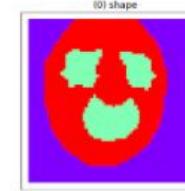
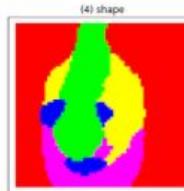
Applications



鵝詞荊陵讌饗貌傷蠶蹤搦襖鄆高憲府
滌咀麯婆芹琬度漁梢顛什僵銳絨莫干
翼唾彈縵磐箋心櫻櫑鼎垂憮睇顆闇摺
塗暉甦娼梟駢丙倫嶠齧哂馨祁廷騰僂
圖暎面苑偷脰鄙抿痕為瘠鱠嶙犧遼鳴
貪鷄駁衲惱庶韻刺貌藁迺旋嫗岱觀灔
躋婢諉甕颺坫汙濶轡赳紊墟兑絅鼯怛
飜狗熾滌保汎之穢捨死鞬窈珥扳壘幃
獨瞻博因眠啟寢併潾膩辯慝虧殉懃櫞
星冠造漆逛鶴僻疋眭嚼睽侈胥險兦崢
榆刃鞚繫怯椹瞞鱸邇鮒卽迺贖敝媾汨
沽鏤磁您蔡忍擣肥悞揠羅卷猷鎗鉦叵
犄肱群堅錳俟痲渝有寘露痂釀忽時挂
鉤踞棘纂鉢擎蠭堋蠻炸瘰鋒疣奕墳椒
灣駢敵操悞顧楨媿壠融寇佛緊鑊封踽
焌麝睨寫葱竣瀝癥哚刺靄坐口妣法鄂

Architecture



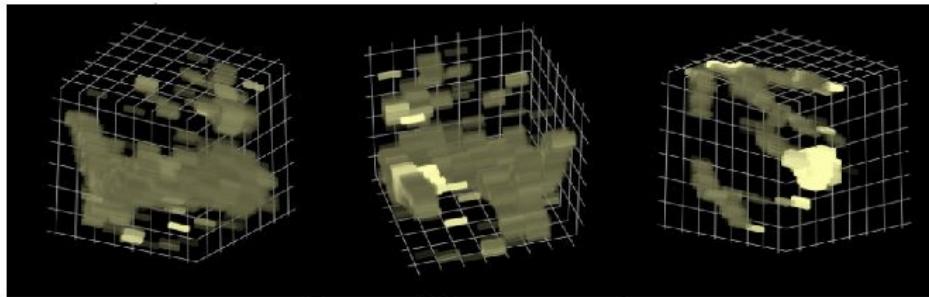


GAN based AOI

- Data
 - 200K, NG and OK are about 1:2.5
 - We use only 50% of data
- StyleGAN2
 - Generate 100K NG fake
 - FID is about 50
- Result(NG as positive)
 - Specificity \geq 70%, FNR 0.018%
 - =train with 100% data(without using GAN)
 - More stable

電腦視覺應用

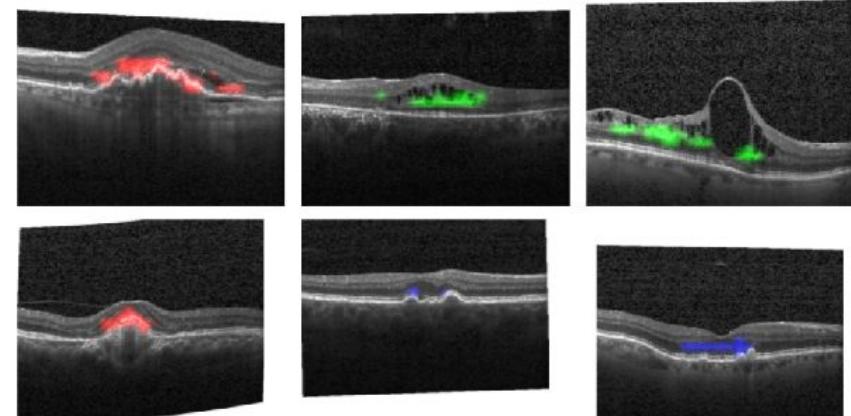
醫療 3D 結核偵測



公安問題偵測



視網膜病變部位偵測



區域偵測範例：斑馬



(複選) MetaGAN 設計用來搭配

